

**Université  
de Limoges**

**UNIVERSITY OF LIMOGES**

Faculty of Science and Technology

Master 1

**Cryptology and Information Security  
(CRYPTIS)**

*Informatic Course*

Network Infrastructure - Semester 2

---

**Implement VLAN trunking with  
IPsec-secured L2TPv3 tunnel**

---

**DUONG Dang-Hung**

**Professor: Pierre-François BONNEFOI**

May 7, 2023

# Table of contents

<b>1</b>	<b>Build the architecture</b>	<b>2</b>
<b>2</b>	<b>Test the functioning of VLANs</b>	<b>7</b>
<b>3</b>	<b>Study L2TPv3 and VXLAN technology</b>	<b>9</b>
3.1	L2TPv3 vs VXLAN . . . . .	9
3.2	Implement VXLAN in Linux with Open vSwitch . . . . .	10
3.3	L2TPv3 traffic encryption solutions . . . . .	10
3.4	L2TPv3/VXLAN vs MPLS . . . . .	11
<b>4</b>	<b>Implementing L2TPv3 tunnel</b>	<b>12</b>
4.1	Verify ARP protocol's functioning . . . . .	13
4.2	Verify the DHCP service . . . . .	14
4.3	Establish TCP connection through the tunnel . . . . .	16
<b>5</b>	<b>Compare different tunnel options</b>	<b>17</b>
5.1	IP encapsulation vs. UDP encapsulation . . . . .	17
5.2	Compare to tunnel made with GRE in GRE-TAP mode . . . . .	19
<b>6</b>	<b>Setting up IPsec encryption on L2TPv3 tunnel</b>	<b>21</b>
<b>7</b>	<b>“Smart” Internet Access</b>	<b>27</b>
7.1	Provide Internet access to the local network . . . . .	27
7.2	Provide Internet access to the hosts . . . . .	28
7.3	Configure default gateway for the hosts . . . . .	29
<b>8</b>	<b>Blocking traffic between VLANs</b>	<b>32</b>
8.1	Using iptables rules . . . . .	32
8.2	Using Routing Policy . . . . .	33

# 1 Build the architecture

From this original network topology, where the local networks (`resD`, `resE`, `resF`, `resG`) are disjoint at layer 2:

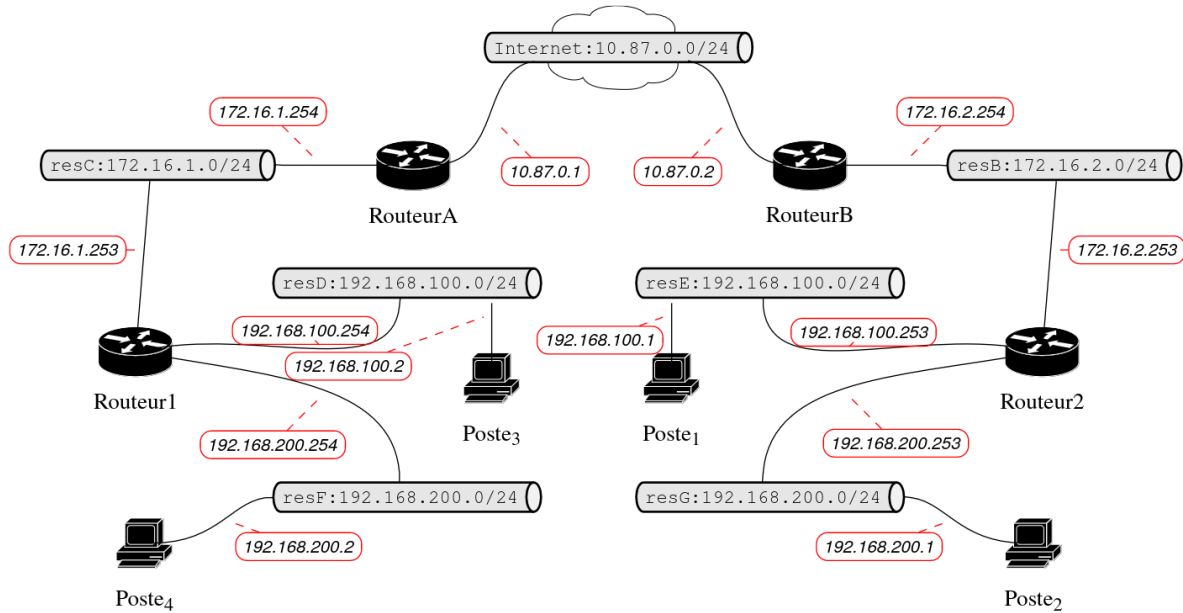


Figure 1: Original network topology

However, our goal is to make the local networks connected at layer 2 and build the following network in Figure 2 below:

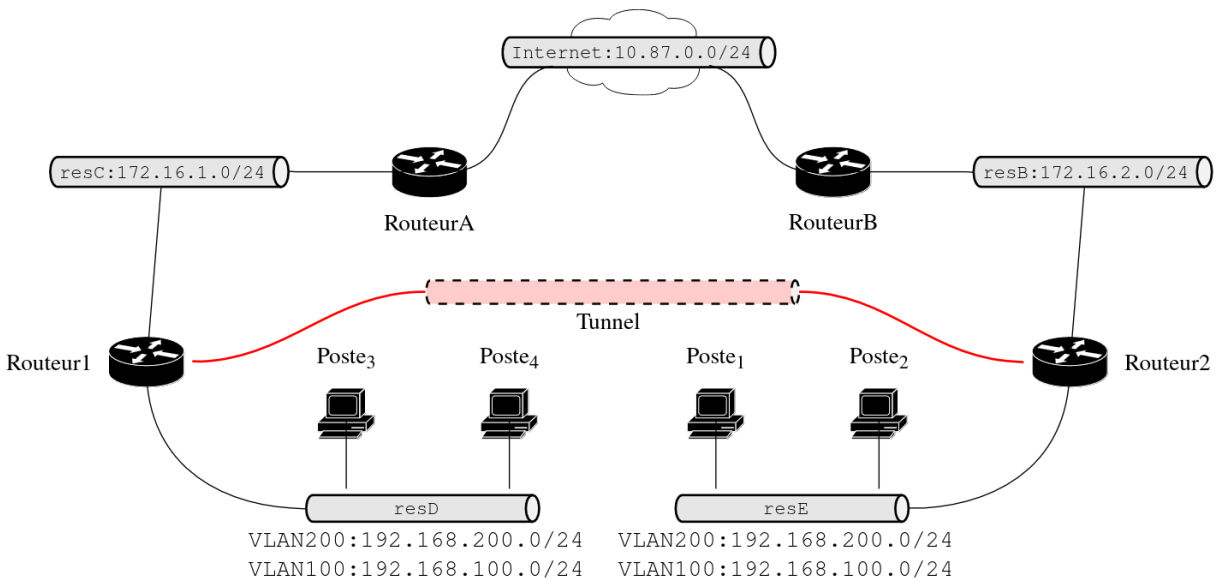


Figure 2: Desired network topology

The project starts with building the architecture. Firstly, we create the net namespaces and switches, corresponding to the routers, work stations and their networks:

```
1 ip netns add routA
2 ip netns add routB
3 ip netns add rout1
4 ip netns add rout2
5 ip netns add poste1
6 ip netns add poste2
7 ip netns add poste3
8 ip netns add poste4
```

```
1 ovs-vsctl add-br internet
2 ovs-vsctl add-br resB
3 ovs-vsctl add-br resC
4 ovs-vsctl add-br resD
5 ovs-vsctl add-br resE
```

Then, we configure each routers and hosts. This includes creating and adding the links between devices, turn the links up, assign them with IP address, configure their routing tables, etc.

Configure routA:

```
1 ip link add routA-eth0 type veth peer name resC-routA
2 ip link add routA-eth1 type veth peer name internet-routA
3 ip link set routA-eth0 netns routA
4 ip link set routA-eth1 netns routA
5 ovs-vsctl add-port resC resC-routA
6 ip link set dev resC-routA up
7 ovs-vsctl add-port internet internet-routA
8 ip link set dev internet-routA up
9 ip netns exec routA ip link set dev lo up
10 ip netns exec routA ip link set dev routA-eth0 up
11 ip netns exec routA ip link set dev routA-eth1 up
12 ip netns exec routA ip a add dev routA-eth0 172.16.1.254/24
13 ip netns exec routA ip a add dev routA-eth1 10.87.0.1/24
14 ip netns exec routA ip r add 172.16.2.0/24 via 10.87.0.2 dev routA-eth1
15 ip netns exec routA sysctl net.ipv4.conf.all.forwarding=1
```

Configure routB:

```
1 ip link add routB-eth0 type veth peer name resB-routB
2 ip link add routB-eth1 type veth peer name internet-routB
3 ip link set routB-eth0 netns routB
4 ip link set routB-eth1 netns routB
5 ovs-vsctl add-port resB resB-routB
6 ip link set dev resB-routB up
7 ovs-vsctl add-port internet internet-routB
8 ip link set dev internet-routB up
9 ip netns exec routB ip link set dev lo up
10 ip netns exec routB ip link set dev routB-eth0 up
11 ip netns exec routB ip link set dev routB-eth1 up
12 ip netns exec routB ip a add dev routB-eth0 172.16.2.254/24
13 ip netns exec routB ip a add dev routB-eth1 10.87.0.2/24
14 ip netns exec routB ip r add 172.16.1.0/24 via 10.87.0.1 dev routB-eth1
15 ip netns exec routB sysctl net.ipv4.conf.all.forwarding=1
```

Configure rout1:

```
1 ip link add rout1-eth0 type veth peer name resC-rout1
2 ip link add rout1-eth1 type veth peer name resD-rout1
3 ip link set rout1-eth0 netns rout1
4 ip link set rout1-eth1 netns rout1
5 ovs-vsctl add-port resC resC-rout1
6 ovs-vsctl add-port resD resD-rout1 trunks=100,200
7 ip link set dev resC-rout1 up
8 ip link set dev resD-rout1 up
9 ip netns exec rout1 ip link set dev lo up
10 ip netns exec rout1 ip link set dev rout1-eth0 up
11 ip netns exec rout1 ip link set dev rout1-eth1 up
12 ip netns exec rout1 ip a add dev rout1-eth0 172.16.1.253/24
13 ip netns exec rout1 ip r add 10.87.0.0/24 via 172.16.1.254 dev rout1-eth0
14 ip netns exec rout1 ip r add 172.16.2.0/24 via 172.16.1.254 dev rout1-eth0
15 ip netns exec rout1 sysctl net.ipv4.conf.all.forwarding=1
```

Configure rout2:

```
1 ip link add rout2-eth0 type veth peer name resB-rout2
2 ip link add rout2-eth1 type veth peer name resE-rout2
3 ip link set rout2-eth0 netns rout2
4 ip link set rout2-eth1 netns rout2
5 ovs-vsctl add-port resB resB-rout2
6 ovs-vsctl add-port resE resE-rout2 trunks=100,200
7 ip link set dev resB-rout2 up
8 ip link set dev resE-rout2 up
9 ip netns exec rout2 ip link set dev lo up
10 ip netns exec rout2 ip link set dev rout2-eth0 up
11 ip netns exec rout2 ip link set dev rout2-eth1 up
12 ip netns exec rout2 ip a add dev rout2-eth0 172.16.2.253/24
13 ip netns exec rout2 ip r add 10.87.0.0/24 via 172.16.2.254 dev rout2-eth0
14 ip netns exec rout2 ip r add 172.16.1.0/24 via 172.16.2.254 dev rout2-eth0
15 ip netns exec rout2 sysctl net.ipv4.conf.all.forwarding=1
```

Configure poste1:

```
1 ip link add poste1-eth0 type veth peer name resE-poste1
2 ip link set poste1-eth0 netns poste1
3 ip netns exec poste1 ip link set dev poste1-eth0 address 00:11:22:00:01:00
4 ovs-vsctl add-port resE resE-poste1 tag=100
5 ip link set dev resE-poste1 up
6 ip netns exec poste1 ip link set dev lo up
7 ip netns exec poste1 ip link set dev poste1-eth0 up
8 ip netns exec poste1 ip a add dev poste1-eth0 192.168.100.1/24
9 ip netns exec poste1 ip r add default dev poste1-eth0 via 192.168.100.254
```

Configure poste2:

```
1 ip link add poste2-eth0 type veth peer name resE-poste2
2 ip link set poste2-eth0 netns poste2
3 ip netns exec poste2 ip link set dev poste2-eth0 address 00:11:22:00:02:00
4 ovs-vsctl add-port resE resE-poste2 tag=200
5 ip link set dev resE-poste2 up
6 ip netns exec poste2 ip link set dev lo up
7 ip netns exec poste2 ip link set dev poste2-eth0 up
8 ip netns exec poste2 ip a add dev poste2-eth0 192.168.200.1/24
9 ip netns exec poste2 ip r add default dev poste2-eth0 via 192.168.200.254
```

Configure poste3:

```
1 ip link add poste3-eth0 type veth peer name resD-poste3
2 ip link set poste3-eth0 netns poste3
3 ip netns exec poste3 ip link set dev poste3-eth0 address 00:11:22:00:03:00
4 ovs-vsctl add-port resD resD-poste3 tag=100
5 ip link set dev resD-poste3 up
6 ip netns exec poste3 ip link set dev lo up
7 ip netns exec poste3 ip link set dev poste3-eth0 up
8 ip netns exec poste3 ip a add dev poste3-eth0 192.168.100.2/24
9 ip netns exec poste3 ip r add default dev poste3-eth0 via 192.168.100.254
```

Configure poste4:

```
1 ip link add poste4-eth0 type veth peer name resD-poste4
2 ip link set poste4-eth0 netns poste4
3 ip netns exec poste4 ip link set dev poste4-eth0 address 00:11:22:00:04:00
4 ovs-vsctl add-port resD resD-poste4 tag=200
5 ip link set dev resD-poste4 up
6 ip netns exec poste4 ip link set dev lo up
7 ip netns exec poste4 ip link set dev poste4-eth0 up
8 ip netns exec poste4 ip a add dev poste4-eth0 192.168.200.2/24
9 ip netns exec poste4 ip r add default dev poste4-eth0 via 192.168.200.254
```

Note that in the configuration of `route1` and `route2`, we created 2 VLANs with ID 100 and 200 to be able to separate the communication inside specific networks 192.168.100.0/24 and 192.168.200.0/24. For this to function, we have to add particular options `trunks` and `tag` when connecting the hosts and routers to the switches, so that the switches are configured to transfer traffic with “VLAN tagging” and “VLAN trunking”:

- The “tag” option is used to configure a port to have a specific VLAN tag (or multiple tags if using the `- tag=VLAN1 - tag=VLAN2` syntax). This means that any traffic received on this port with the specified VLAN tag(s) will be associated with the specified VLAN(s) and forwarded accordingly. Thus we apply it on the connection link between the routers and the switches.
- The “trunks” option is used to configure a port as a trunk port, which means that it will carry traffic from multiple VLANs. Any traffic received on this port will not have a specific VLAN tag associated with it, but will rather have a VLAN tag added by the switch based on the VLAN configuration. Thus we apply it on the connection link between the hosts (postes) and the switches.

Also, for the configuration to be more convenient in the following exercises, I decided to set up static MAC addresses for the network cards of the host devices (which is quite reasonable since in real life each device is also assigned with a constant MAC address instead of having random ones):

- poste1-eth0: 00:11:22:00:01:00
- poste2-eth0: 00:11:22:00:02:00
- poste3-eth0: 00:11:22:00:03:00
- poste4-eth0: 00:11:22:00:04:00

After having finished constructing the network topology, we would have this graph (since the network is quite big, the graph might need to be zoomed to be seen):


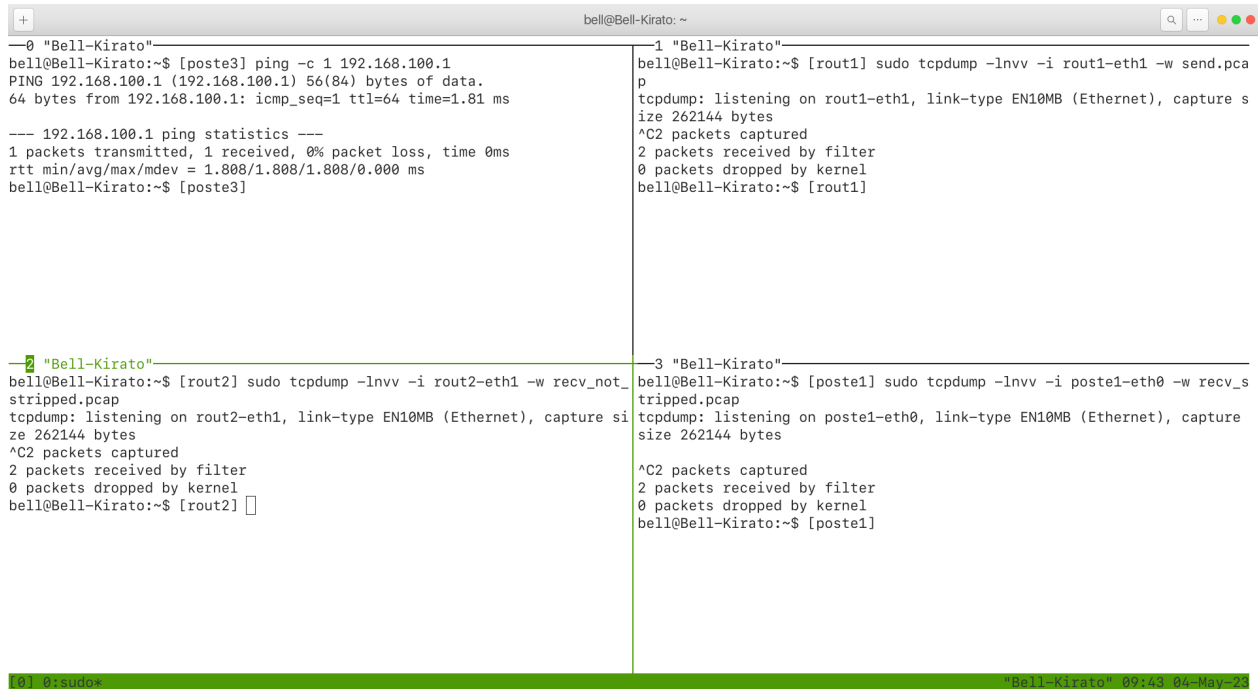


Figure 3: Network graph

## 2 Test the functioning of VLANs

We will try to test the functioning of VLANs configured and inspect the nature of the traffic to verify the functioning of “VLAN tagging”. In this example, we will ping from **poste1** to **poste3** (both in VLAN100) (the process is the same in the case of VLAN200), while sniffing on 2 VLAN trunking interfaces **rout1-eth1**, **rout1-eth2**, and the interface of the destination host **poste1-eth0** using **tcpdump**:



```
bell@Bell-Kirato: ~  
--0 "Bell-Kirato"  
bell@Bell-Kirato:~$ [poste3] ping -c 1 192.168.100.1  
PING 192.168.100.1 (192.168.100.1) 56(84) bytes of data.  
64 bytes from 192.168.100.1: icmp_seq=1 ttl=64 time=1.81 ms  
  
--- 192.168.100.1 ping statistics ---  
1 packets transmitted, 1 received, 0% packet loss, time 0ms  
rtt min/avg/max/mdev = 1.808/1.808/1.808/0.000 ms  
bell@Bell-Kirato:~$ [poste3]  
  
--2 "Bell-Kirato"  
bell@Bell-Kirato:~$ [rout2] sudo tcpdump -lnvv -i rout2-eth1 -w recv_not_stripped.pcap  
tcpdump: listening on rout2-eth1, link-type EN10MB (Ethernet), capture size 262144 bytes  
^C2 packets captured  
2 packets received by filter  
0 packets dropped by kernel  
bell@Bell-Kirato:~$ [rout2]   
  
--1 "Bell-Kirato"  
bell@Bell-Kirato:~$ [rout1] sudo tcpdump -lnvv -i rout1-eth1 -w send.pcap  
tcpdump: listening on rout1-eth1, link-type EN10MB (Ethernet), capture size 262144 bytes  
^C2 packets captured  
2 packets received by filter  
0 packets dropped by kernel  
bell@Bell-Kirato:~$ [rout1]  
  
--3 "Bell-Kirato"  
bell@Bell-Kirato:~$ [poste1] sudo tcpdump -lnvv -i poste1-eth0 -w recv_stripped.pcap  
tcpdump: listening on poste1-eth0, link-type EN10MB (Ethernet), capture size 262144 bytes  
^C2 packets captured  
2 packets received by filter  
0 packets dropped by kernel  
bell@Bell-Kirato:~$ [poste1]
```

Figure 4: Ping from **poste3** to **poste1**

Here, the packets captured are saved in **.pcap** files so as to be analyzed later:

- Packets received on **rout1-eth1** are saved in **send.pcap**
- Packets received on **rout2-eth1** are saved in **recv\_not\_stripped.pcap**
- Packets received on **poste1-eth0** are saved in **recv\_stripped.pcap**

Note that we have configured the MAC address of **poste1** to be **00:11:22:00:01:00**, and that of **poste3** to be **00:11:22:00:03:00**. This information will be shown in the following Figure 5.



To inspect these captured packets, **scapy** is utilized:

```

>>> send = rdpcap("send.pcap")[0]
>>> send.show()
### [ Ethernet ]###
  dst      = 00:11:22:00:01:00
  src      = 00:11:22:00:03:00
  type     = n_802_1Q
### [ 802.1Q ]###
  prio     = 0
  id       = 0
  vlan     = 100
  type     = IPv4
### [ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 84
  id       = 27349
  flags    = DF
  frag     = 0
  ttl      = 64
  proto    = icmp
  chksum   = 0x867f
  src      = 192.168.100.2
  dst      = 192.168.100.1
  \options \
### [ ICMP ]###
  type     = echo-request
  code     = 0
  chksum   = 0x9281
  id       = 0x602d
  seq      = 0x1
  unused   = ''
### [ Raw ]###
  load     = '\x16bSd\x00\x00\x00\x00\x00\x00\x00\x00\x10\x11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f4\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f'

[1] 0:python3+

...: pcap")[0]
>>> recv_not_stripped = rdpcap("recv_not_stripped.pcap")[0]
>>> recv_not_stripped.show()
### [ Ethernet ]###
  dst      = 00:11:22:00:01:00
  src      = 00:11:22:00:03:00
  type     = n_802_1Q
### [ 802.1Q ]###
  prio     = 0
  id       = 0
  vlan     = 100
  type     = IPv4
### [ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 84
  id       = 27349
  flags    = DF
  frag     = 0
  ttl      = 64
  proto    = icmp
  chksum   = 0x867f
  src      = 192.168.100.2
  dst      = 192.168.100.1
  \options \
### [ ICMP ]###
  type     = echo-request
  code     = 0
  chksum   = 0x9281
  id       = 0x602d
  seq      = 0x1
  unused   = ''
### [ Raw ]###
  load     = '\x16bSd\x00\x00\x00\x00\x00\x00\x00\x00'

...: [0]
>>> recv_stripped = rdpcap("recv_stripped.pcap")
>>> recv_stripped.show()
### [ Ethernet ]###
  dst      = 00:11:22:00:01:00
  src      = 00:11:22:00:03:00
  type     = IPv4
### [ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 84
  id       = 27349
  flags    = DF
  frag     = 0
  ttl      = 64
  proto    = icmp
  chksum   = 0x867f
  src      = 192.168.100.2
  dst      = 192.168.100.1
  \options \
### [ ICMP ]###
  type     = echo-request
  code     = 0
  chksum   = 0x9281
  id       = 0x602d
  seq      = 0x1
  unused   = ''
### [ Raw ]###
  load     = '\x16bSd\x00\x00\x00\x00\x00\x00\x00\x00\x10\x11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !"#%&'()*+,-./01234567'

>>>

```

Figure 5: Analyzing packets handled by VLAN tagging with **scapy**

From Figure 5, we see that the packets, which are the same ICMP Request from **poste3**, received on interface **route1-eth1** and **route2-eth1** has an additional 802.1Q header, with a field **vlan = 100**. This is exactly the VLAN header (or the *VLAN tag*) which we are expecting. The process can be explained as follows:

- **poste3** sends an ICMP Request to the destination 192.168.100.1 (**poste1**)
- The packet travels to the port **resD-poste3** of switch **resD** which was configured with **tag=100**
- The packet arrives at the port and is injected with the 802.1Q header which has a tag **vlan = 100**, then it is forwarded to **route1**
- The packet goes from **route1** to **route2** (through the tunnel, which we will discuss in the following sessions)
- The packet arrives at **route2** and continues to switch **resE** using interface **route2-eth1**
- The packet travels to the port **resE-route2** which was configured with **trunks=100,200**
- Since the packet was tagged with the VLAN ID: **vlan = 100**, which is included in the option **trunks=100,200**, switch **resE** finds the port configured with **tag=100**, which is **resE-poste1**, or the one peered with **poste1-eth0**, strips the VLAN tag, and then forward the packet to that output port
- **poste1** receives the packet with no VLAN tag

## 3 Study L2TPv3 and VXLAN technology

### 3.1 L2TPv3 vs VXLAN

To compare L2TPv3 and VXLAN, we will briefly study their similarities and differences.

Overall, they both encapsulate layer 2 packets inside IP packets and transport them over an IP network, for example the Internet. However, apart from this point, there are a couple of differences between their purposes and applications.

Regarding L2TPv3:

- It is an improved version of L2TP (Layer 2 Tunneling Protocol), which is a tunneling protocol.
- As a tunnel, it has 2 endpoints only: L2TP access concentrator (LAC) and the L2TP network server (LNS).
- It allows the creation of a virtual private dialup network (VPDN) to connect a remote client to its corporate network by using a shared infrastructure, which could be the Internet or a service provider's network.
- It is a layer 2 protocol, which means it was designed to provide a way to transport data link layer protocols, such as Ethernet frames, over an IP network.

On the other hand, VXLAN:

- Is a network virtualization technology.
- Is capable of creating an overlay network or tunnel.
- Can have multiple endpoints named VTEPs (VXLAN Tunnel EndPoints)
- It provides a scalable and flexible method for building virtual networks that can span multiple physical networks → VXLAN plays an important role in cloud and data center environments.
- Is a protocol used to extend Layer 2 networks over a Layer 3 infrastructure.
- Is typically used in data center environments to provide network virtualization and overlay networks.
- It operates between layer 2 and layer 3 → A “layer 2.5” protocol.

## 3.2 Implement VXLAN in Linux with Open vSwitch

In order to implement VXLAN in Linux with Open vSwitch, it is essential to create a bridge with type VXLAN, configure the bridge as well as the system appropriately to connect to the VXLAN. For example, in our project, if we would like to create a VXLAN for VLAN100: On `poste1`:

- Create a bridge:

```
1 ovs-vsctl add-br br0
```

- Turn the bridge up:

```
1 ip link set br0 up
```

- Create the VXLAN interface to represent its VTEP:

```
1 ovs-vsctl add-port br0 vx_poste3 -- set interface vx_poste3 type=vxlan
  options:remote_ip=192.168.100.2 options:local_ip=192.168.100.1
```

Symmetrically, we configure `poste3`:

- Create a bridge:

```
1 ovs-vsctl add-br br0
```

- Turn the bridge up:

```
1 ip link set br0 up
```

- Create the VXLAN interface to represent its VTEP:

```
1 ovs-vsctl add-port br0 vx_poste1 -- set interface vx_poste1 type=vxlan
  options:remote_ip=192.168.100.1 options:local_ip=192.168.100.2
```

The overlay network is now operational, and everything connected to `br0` on either machine will appear to be connected to the same layer 2 network.

## 3.3 L2TPv3 traffic encryption solutions

Compared to the original L2TP, L2TPv3 supports encryption of tunnel traffic using a variety of encryption algorithms, such as 3DES, AES, and Blowfish, therefore provides improved security through the use of IPsec encryption and authentication, which can be used to protect the L2TPv3 control and data packets. When IPsec is used for encryption, it provides security at the network layer and creates a virtual private network (VPN) tunnel between the L2TPv3 tunnel endpoints. This enables secure transmission of data over the public Internet. This type of combination is well-known as “L2TPv3 over IPsec”, which we might see later. Besides IPsec, there are several other options such as SSL/TLS, MACsec, AAA, etc.

### 3.4 L2TPv3/VXLAN vs MPLS

In comparison, they have a few similarities:

- All of them allows for the creation of VPNs (Virtual Private Networks) by using labels to segregate traffic from different VPNs, offering secure and scalable connectivity between geographically dispersed sites.
- They make routing decisions for packets by labelling the packet with an additional header.
- All three protocols support encryption and other security mechanisms to protect data in transit.

However, there are some significant differences:

- L2TPv3 is layer 2 protocol, VLAN and MPLS are layer 2.5 technologies.
- Encapsulation:
  - L2TPv3 and VXLAN: encapsulate Layer 2 frames within IP packets
  - MPLS encapsulates IP packets within MPLS headers
- Main purposes:
  - MPLS is designed to improve the performance and efficiency of packet forwarding by using labels to route packets through the network (used to transport layer 3 packets).
  - L2TPv3 is designed to provide a way to transport data link layer protocols over an IP network (used to transport layer 2 packets).
  - VXLAN is designed to building overlaying virtual networks that can span multiple physical networks (used to transport layer 2 packets).
- Along the path of the packet:
  - MPLS: all immediate routers have to examine the label to determine the next hop, the final router in the path removes the label and forwards the packet to its destination → MPLS uses label-switching to forward traffic.
  - L2TPv3 and VXLAN: only the endpoints (tunnel endpoints in the case of L2TPv3, VTEPs in the case of VXLAN) have to look at the label, remove it and forwards the packet to its destination → L2TPv3 and VXLAN use traditional routing protocols.
- To provide end-to-end connectivity for IP-based services:
  - L2TPv3 and VXLAN are widely used within data center and enterprise networks
  - MPLS is usually used within service provider networks

## 4 Implementing L2TPv3 tunnel

The process of creating the L2TPv3 tunnel in IP encapsulation mode was very well-guided in the subject document.

We implement the tunnel on rout1 side:

```
1 # Config tunnel rout1
2 ip netns exec rout1 ip l2tp add tunnel remote 172.16.2.253 local
   172.16.1.253 encap ip tunnel_id 3000 peer_tunnel_id 4000
3 ip netns exec rout1 ip l2tp add session tunnel_id 3000 session_id 1000
   peer_session_id 2000
4 ip netns exec rout1 ip link set l2tpeth0 up
5 ip netns exec rout1 ip a flush dev rout1-eth1
6 ip netns exec rout1 brctl addbr tunnel
7 ip netns exec rout1 brctl addif tunnel l2tpeth0
8 ip netns exec rout1 brctl addif tunnel rout1-eth1
9 ip netns exec rout1 ip link set tunnel up
10 ip netns exec rout1 ip link add link tunnel name tunnel.100 type vlan id
   100
11 ip netns exec rout1 ip link set tunnel.100 up
12 ip netns exec rout1 ip link add link tunnel name tunnel.200 type vlan id
   200
13 ip netns exec rout1 ip link set tunnel.200 up
14 ip netns exec rout1 ip addr add 192.168.100.254/24 dev tunnel.100
15 ip netns exec rout1 ip addr add 192.168.200.254/24 dev tunnel.200
16 ip netns exec rout1 ip link set dev l2tpeth0 mtu 1500
```

Symmetrically, we implement the tunnel on rout2 side:

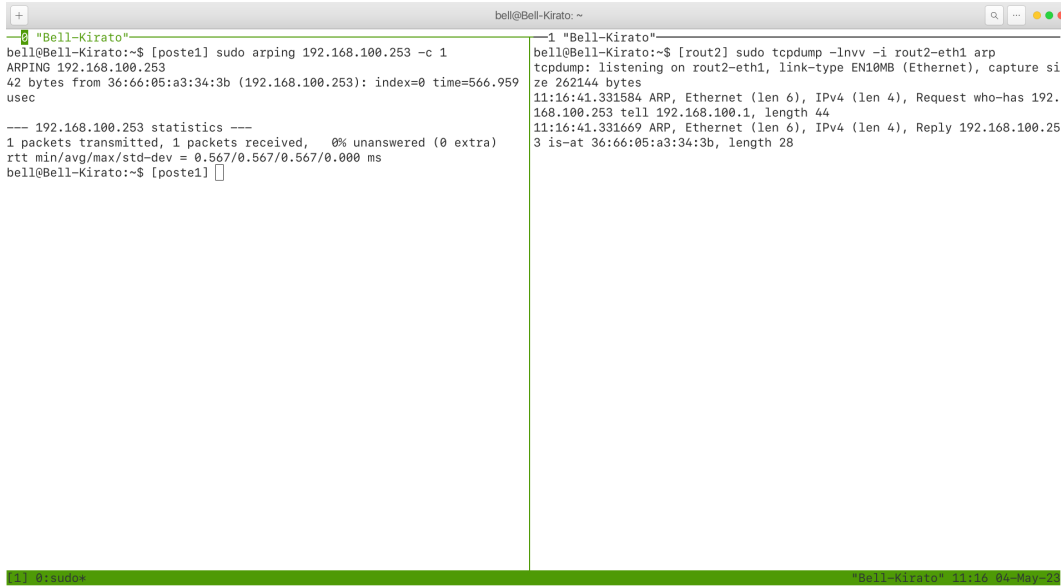
```
1 # Config tunnel rout2
2 ip netns exec rout2 ip l2tp add tunnel remote 172.16.1.253 local
   172.16.2.253 encap ip tunnel_id 4000 peer_tunnel_id 3000
3 ip netns exec rout2 ip l2tp add session tunnel_id 4000 session_id 2000
   peer_session_id 1000
4 ip netns exec rout2 ip link set l2tpeth0 up
5 ip netns exec rout2 ip a flush dev rout2-eth1
6 ip netns exec rout2 brctl addbr tunnel
7 ip netns exec rout2 brctl addif tunnel l2tpeth0
8 ip netns exec rout2 brctl addif tunnel rout2-eth1
9 ip netns exec rout2 ip link set tunnel up
10 ip netns exec rout2 ip link add link tunnel name tunnel.100 type vlan id
   100
11 ip netns exec rout2 ip link set tunnel.100 up
12 ip netns exec rout2 ip link add link tunnel name tunnel.200 type vlan id
   200
13 ip netns exec rout2 ip link set tunnel.200 up
14 ip netns exec rout2 ip addr add 192.168.100.253/24 dev tunnel.100
15 ip netns exec rout2 ip addr add 192.168.200.253/24 dev tunnel.200
16 ip netns exec rout2 ip link set l2tpeth0 mtu 1500
```

Now that the tunnel has been configured, we can try to see if it was installed successfully and functions well.

## 4.1 Verify ARP protocol's functioning

First, we will verify to see if the ARP protocol works normally for the discovery of machines regardless of their side connection to the Internet.

In this example, we will try to send ARP Requests from **poste1** to **rou1** (same side) and **poste3** (opposite side, same VLAN) using the **arping** command.



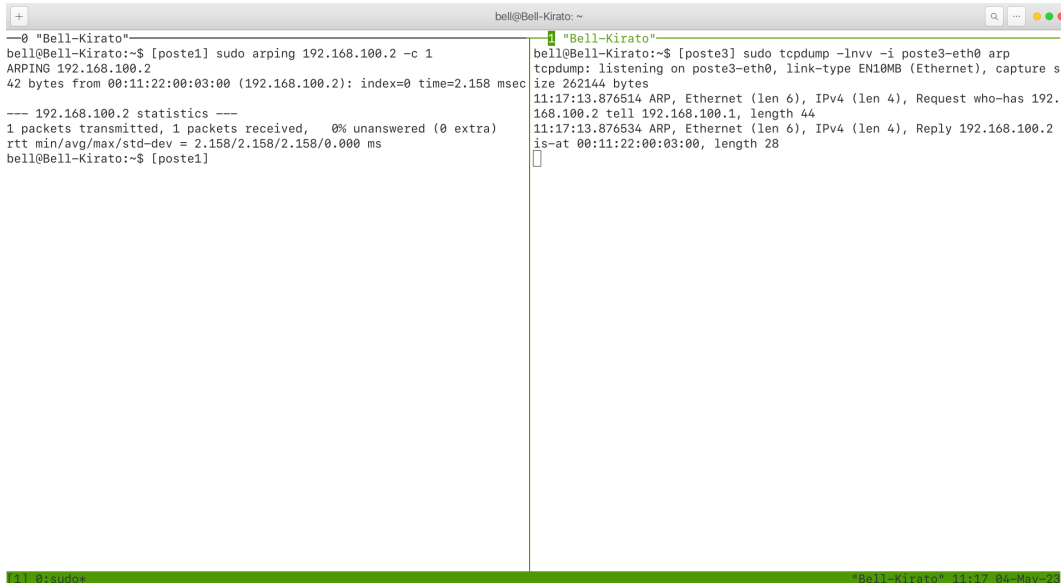
The screenshot shows a terminal window with two panes. The left pane shows the execution of the `arping` command from `poste1` to `192.168.100.253`. The output indicates a successful ARP request with a response time of 566.959 microseconds. The right pane shows the output of `tcpdump` on `rou1`, capturing the ARP request and reply packets between `192.168.100.1` and `192.168.100.253`.

```
bell@Bell-Kirato:~$ [poste1] sudo arping 192.168.100.253 -c 1
ARPING 192.168.100.253
42 bytes from 36:66:05:a3:34:3b (192.168.100.253): index=0 time=566.959
usec

--- 192.168.100.253 statistics ---
1 packets transmitted, 1 packets received, 0% unanswered (0 extra)
rtt min/avg/max/std-dev = 0.567/0.567/0.567/0.000 ms
bell@Bell-Kirato:~$ [poste1]
```

```
bell@Bell-Kirato:~$ [rou1] sudo tcpdump -lnvv -i rout2-eth1 arp
tcpdump: listening on rout2-eth1, link-type EN10MB (Ethernet), capture si
ze 262144 bytes
11:16:41.331584 ARP, Ethernet (len 6), IPv4 (len 4), Request who-has 192.
168.100.253 tell 192.168.100.1, length 44
11:16:41.331669 ARP, Ethernet (len 6), IPv4 (len 4), Reply 192.168.100.25
3 is-at 36:66:05:a3:34:3b, length 28
```

Figure 6: ARP Protocol between **poste1** and **rou1** (same side)



The screenshot shows a terminal window with two panes. The left pane shows the execution of the `arping` command from `poste1` to `192.168.100.2`. The output indicates a successful ARP request with a response time of 2.158 milliseconds. The right pane shows the output of `tcpdump` on `poste3`, capturing the ARP request and reply packets between `192.168.100.1` and `192.168.100.2`.

```
bell@Bell-Kirato:~$ [poste1] sudo arping 192.168.100.2 -c 1
ARPING 192.168.100.2
42 bytes from 00:11:22:00:03:00 (192.168.100.2): index=0 time=2.158 msec

--- 192.168.100.2 statistics ---
1 packets transmitted, 1 packets received, 0% unanswered (0 extra)
rtt min/avg/max/std-dev = 2.158/2.158/2.158/0.000 ms
bell@Bell-Kirato:~$ [poste1]
```

```
bell@Bell-Kirato:~$ [poste3] sudo tcpdump -lnvv -i poste3-eth0 arp
tcpdump: listening on poste3-eth0, link-type EN10MB (Ethernet), capture s
ize 262144 bytes
11:17:13.876514 ARP, Ethernet (len 6), IPv4 (len 4), Request who-has 192.
168.100.2 tell 192.168.100.1, length 44
11:17:13.876534 ARP, Ethernet (len 6), IPv4 (len 4), Reply 192.168.100.2
is-at 00:11:22:00:03:00, length 28
```

Figure 7: ARP Protocol between **poste1** and **poste3** (opposite side)

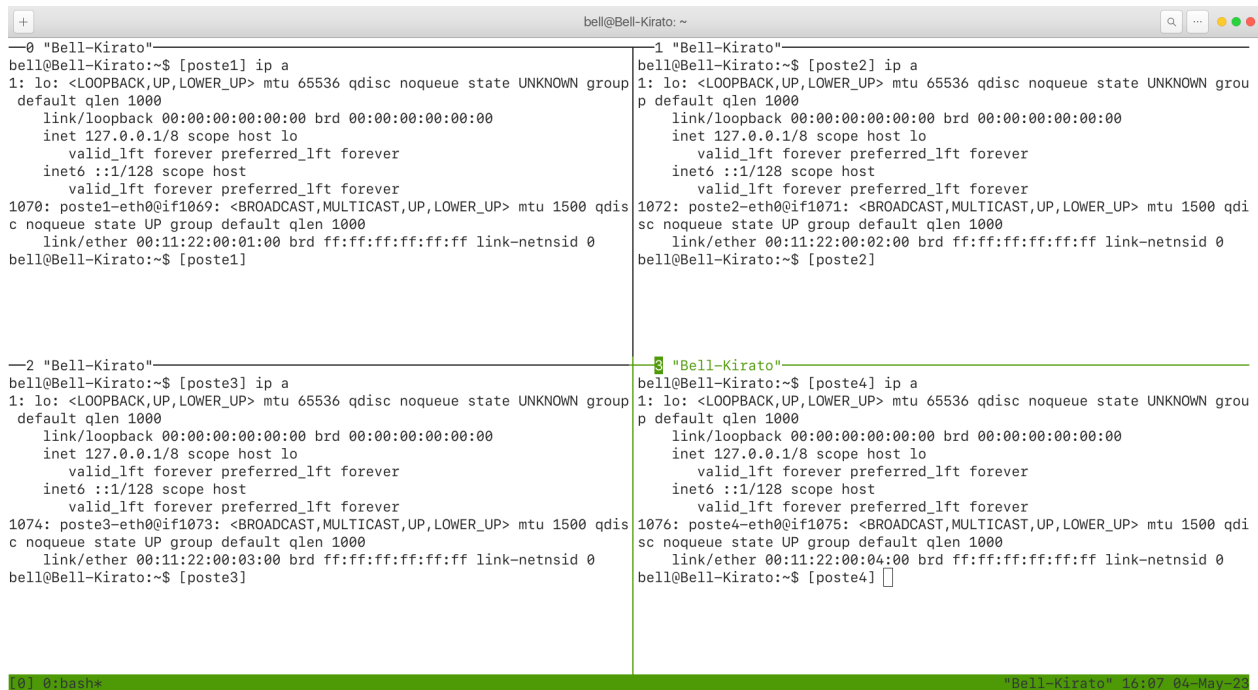
It can be seen from Figure 6 and Figure 8 that the ARP Protocol works well in this case.

## 4.2 Verify the DHCP service

We will implement a DHCP server on `router1` and verify that the DHCP service can be offered to all 4 hosts regardless of their side connection to the internet. Firstly, all the pre-registered IP addresses of the hosts should be flushed:

```
1 $ sudo ip netns exec poste1 ip a flush dev poste1-eth0
2 $ sudo ip netns exec poste2 ip a flush dev poste2-eth0
3 $ sudo ip netns exec poste3 ip a flush dev poste3-eth0
4 $ sudo ip netns exec poste4 ip a flush dev poste4-eth0
```

We check that all the IP addresses of the network interfaces of the hosts are deleted:



```
bell@Bell-Kirato:~$ [poste1] ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group
default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
1070: poste1-eth0@if1069: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc
noqueue state UP group default qlen 1000
    link/ether 00:11:22:00:01:00 brd ff:ff:ff:ff:ff:ff link-netnsid 0
bell@Bell-Kirato:~$ [poste1]

bell@Bell-Kirato:~$ [poste2] ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group
default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
1072: poste2-eth0@if1071: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc
noqueue state UP group default qlen 1000
    link/ether 00:11:22:00:02:00 brd ff:ff:ff:ff:ff:ff link-netnsid 0
bell@Bell-Kirato:~$ [poste2]

bell@Bell-Kirato:~$ [poste3] ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group
default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
1074: poste3-eth0@if1073: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc
noqueue state UP group default qlen 1000
    link/ether 00:11:22:00:03:00 brd ff:ff:ff:ff:ff:ff link-netnsid 0
bell@Bell-Kirato:~$ [poste3]

bell@Bell-Kirato:~$ [poste4] ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group
default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
1076: poste4-eth0@if1075: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc
noqueue state UP group default qlen 1000
    link/ether 00:11:22:00:04:00 brd ff:ff:ff:ff:ff:ff link-netnsid 0
bell@Bell-Kirato:~$ [poste4]
```

Figure 8: Flush all network interfaces of `poste1`, `poste2`, `poste3`, and `poste4`

To configure the DHCP server on `router1`, we configure the file `/etc/dnsmasq.conf` by adding these lines:

```
1 dhcp-range=192.168.100.1,192.168.100.2,12h
2 dhcp-range=192.168.200.1,192.168.200.2,12h
3 dhcp-host=00:11:22:00:01:00,192.168.100.1,12h
4 dhcp-host=00:11:22:00:02:00,192.168.200.1,12h
5 dhcp-host=00:11:22:00:03:00,192.168.100.2,12h
6 dhcp-host=00:11:22:00:04:00,192.168.200.2,12h
```

We used option `dhcp-host` in order to configure static IP addresses for the hosts using their MAC address, and option `dhcp-range` to specify the possible IP address range to assign. This is also the reason why I configure static IP addresses for the hosts at the beginning.

After that, we start the DHCP Server on `router1` using `dnsmasq` command:

```
1 $ sudo ip netns exec rout1 dnsmasq -d -z
2 dnsmasq: started, version 2.80 cachesize 150
3 dnsmasq: compile time options: IPv6 GNU-getopt DBus i18n IDN DHCP DHCPv6
   no-Lua TFTP conntrack ipset auth nettlehash DNSSEC loop-detect inotify
   dumpfile
4 dnsmasq-dhcp: DHCP, IP range 192.168.200.1 -- 192.168.200.2, lease time 12
   h
5 dnsmasq-dhcp: DHCP, IP range 192.168.100.1 -- 192.168.100.2, lease time 12
   h
6 dnsmasq: reading /etc/resolv.conf
7 dnsmasq: using nameserver 127.0.0.53#53
8 dnsmasq: read /etc/hosts - 7 addresses
```

On each host, we request IP address using `dhclient`:

```
1 $ [poste1] sudo dhclient -v poste1-eth0
2 $ [poste2] sudo dhclient -v poste2-eth0
3 $ [poste3] sudo dhclient -v poste3-eth0
4 $ [poste4] sudo dhclient -v poste4-eth0
```

We can confirm that the DHCP Protocol, including DHCP DISCOVER, DHCP OFFER, DHCP REQUEST, DHCP PACK messages, are carried out normally for all 4 hosts, regardless of their side connection to the Internet:

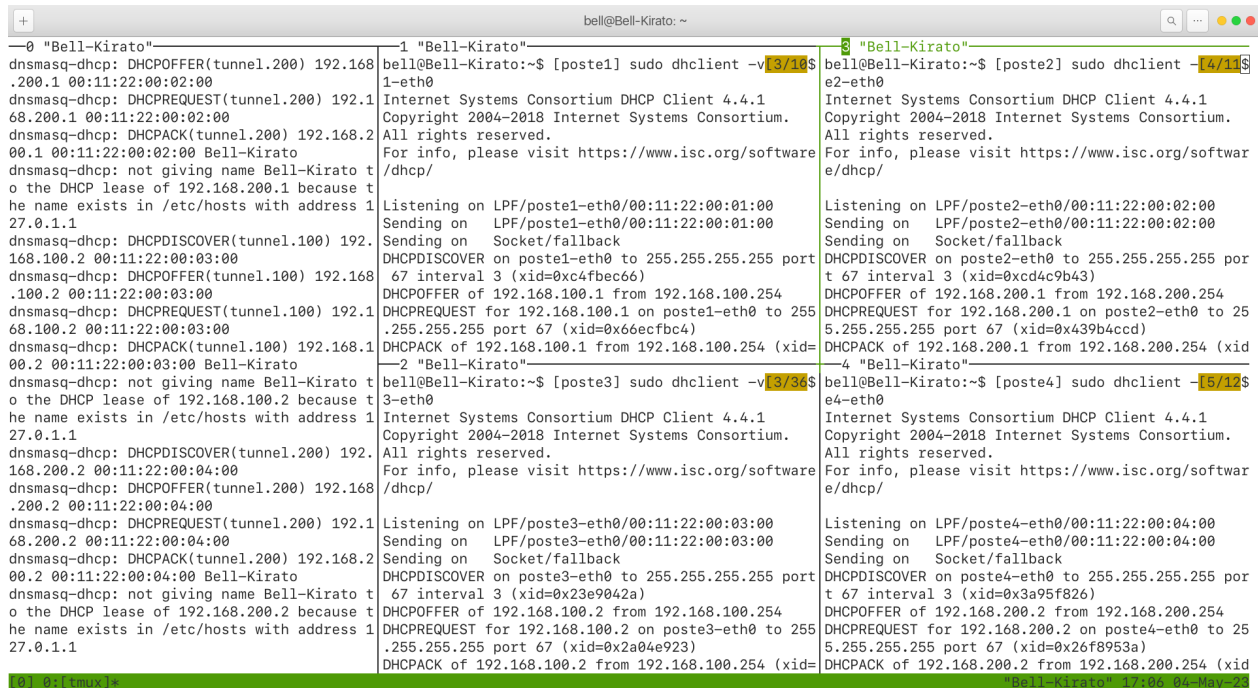


Figure 9: The DHCP Protocol captured on rout1 and poste1, poste2, poste3, and poste4



### 4.3 Establish TCP connection through the tunnel

We also can verify the functioning of the tunnel by establishing a TCP connection between `route1` and `poste1` using `socat`:

```
bell@Bell-Kirato: ~
0 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste1] socat - tcp:192.168.100.254:8080
hello
hi

1 "Bell-Kirato"
bell@Bell-Kirato:~$ [route1] socat - tcp-listen:8080
000
hello
hi

2 "Bell-Kirato"
bell@Bell-Kirato:~$ [route2] sudo tcpdump -lnvX -i route2-eth0
tcpdump: listening on route2-eth0, link-type EN10MB (Ethernet), capture size 262144 bytes
17:27:01.900602 IP (tos 0x0, ttl 64, id 31513, offset 0, flags [none], proto unknown (115), length 104)
    172.16.2.253 > 172.16.1.253: ip-proto-115 84
    0x0000: 4500 0068 7b19 0000 4073 a1ef ac10
02fd E..h[...].s.....
    0x0010: ac10 01fd 0000 03e8 0000 0000 1ad9
9ccd .....
    0x0020: 801a 0011 2200 0100 8100 0064 0800
4500 .....d..E.
    0x0030: 003a 73fd 4000 4006 7c70 c0a8 6401
c0a8 ..s.@.@.|p..d...
$ "Bell-Kirato"
bell@Bell-Kirato:~$ [route2] sudo tcpdump -lnvX -i tunnel.100 tcp
tcpdump: listening on tunnel.100, link-type EN10MB (Ethernet), capture size 262144 bytes
17:26:25.215694 IP (tos 0x0, ttl 64, id 29691, offset 0, flags [DF], proto TCP (6), length 60)
    192.168.100.1.55980 > 192.168.100.254.8080: Flags [S], cksum 0x246b (correct), seq 3059609292, win 64240, options [mss 1460,sackOK,TS val 823528918 ec r 0,nop,wscale 7], length 0
    0x0000: 4500 003c 73fb 4000 4006 7c70 c0a8
6401 E..<s.@.@.|p..d.
    0x0010: c0a8 64fe daac 1f90 b65d eec0 0000
0000 ..d.....].....
    0x0020: a002 faf0 246b 0000 0204 05b4 0402
080a ....$k.....
Bell-Kirato 17:27 04-May-21
```

Figure 10: TCP connection established between `route1` and `poste1`

The TCP connection was established successfully between `route1` and `poste1`. In Figure 13, it can be seen that when we capture the packets on `route2`, the TCP packets arrives on interface `tunnel.100` in their original form (TCP packets), while the packets arrives on interface `route2-eth0` were encapsulated using L2TP protocol (`ip-proto-115`).

## 5 Compare different tunnel options

### 5.1 IP encapsulation vs. UDP encapsulation

Here we would try to compare 2 encapsulation modes for the tunnel: IP and UDP.

IP encapsulation mode is originally applied:

```
1 ip netns exec rout1 ip l2tp add tunnel remote 172.16.2.253 local
  172.16.1.253 encap ip tunnel_id 3000 peer_tunnel_id 4000
2 ip netns exec rout2 ip l2tp add tunnel remote 172.16.1.253 local
  172.16.2.253 encap ip tunnel_id 4000 peer_tunnel_id 3000
```

After the network architecture is built successfully, we listen on interface `routA-eth0` of `routA`, so as to capture the packet encapsulated in IP mode and save them in the file `encap_ip.pcap`:

```
1 $ sudo ip netns exec routA tcpdump -lnvX -i routA-eth0 -w encap_ip.pcap
```

Then we send a ping from `poste3` to `poste1`:

```
1 $ sudo ip netns exec poste3 ping 192.168.100.1 -c 1
```

To use UDP encapsulation mode, we have to change the `encap ip` argument to `encap udp`, and we also need to specify the UDP ports which the routers on each end of the tunnel use to communicate. In this example, we will choose port 8888 for `rout1` and port 9999 for `rout2`:

```
1 ip netns exec rout1 ip l2tp add tunnel remote 172.16.2.253 local
  172.16.1.253 encap udp udp_sport 8888 udp_dport 9999 tunnel_id 3000
  peer_tunnel_id 4000
2 ip netns exec rout2 ip l2tp add tunnel remote 172.16.1.253 local
  172.16.2.253 encap udp udp_sport 9999 udp_dport 8888 tunnel_id 4000
  peer_tunnel_id 3000
```

After the network architecture is re-built successfully, we do the same process as above to capture the packet encapsulated in UDP mode and save them in the file `encap_udp.pcap`:

```
1 $ sudo ip netns exec routA tcpdump -lnvX -i routA-eth0 -w encap_udp.pcap
```

Then once again, we send a ping from `poste3` to `poste1`:

```
1 $ sudo ip netns exec poste3 ping 192.168.100.1 -c 1
```

```
+
bell@Bell-Kirato ~
>>> ip_p = rdpcap('encap_ip.pcap')
>>> ip_p[0].show()
#### Ethernet ####
dst      = c2:74:31:03:e6:8a
src       = ae:8f:79:84:e4:fd
type      = IPv4
#### IP ####
version   = 4
ihl       = 5
tos       = 0x0
len       = 130
id        = 11694
flags     = 
frag      = 0
ttl       = 64
proto     = 12tp
chksum    = 0xef40
src       = 172.16.1.253
dst       = 172.16.2.253
\options
#### Raw ####
load      = '\x00\x00\x07\x0d\x00\x00\x00\x00\x00\x11'\x00\x01\x
00\x00\x11'\x00\x03\x00\x18\x00\x0d\x08\x00E\x00\x00T'\x82%\x00@\x01/\
\x0c\x18a8'd\x02|\xc0\x18a8'd\x01\x08\x00\x1df?|\xbab6\x00\x01\x06|\xfadSd\x00
\x00\x00\x00:\x0b\x00\x00\x00\x00\x00\x10\x11\x12\x13\x14\x15\x16\x17\x1
8\x19\x1a\x1b\x1c\x1d\x1e\x1f !"#$%&'()*+,-./01234567'

>>>
```

```
+
bell@Bell-Kirato ~
>>> udp_p = rdpcap('encap_udp.pcap')
>>> udp_p[1].show()
#### Ethernet ####
dst      = 8a:22:ba:c8:3c:f0
src       = fa:18:60:e8:09:0d
type      = IPv4
#### IP ####
version   = 4
ihl       = 5
tos       = 0x0
len       = 142
id        = 11608
flags     = 
frag      = 0
ttl       = 64
proto     = udp
chksum    = 0xefec
src       = 172.16.1.253
dst       = 172.16.2.253
\options
#### UDP ####
sport     = 8888
dport     = 9999
len       = 122
chksum    = 0x0
#### Raw ####
load      = '\x00\x03\x00\x00\x00\x00\x07\x0d\x00\x00\x00\x00\x00\x00\x00\x00\x11'\x00\x01\x00\x00\x00\x11'\x00\x03\x00\x18\x00\x0d\x08\x00E\x00\x00T'\x82%\x00@\x01/\
\x0c\x18a8'd\x02|\xc0\x18a8'd\x01\x08\x00\x1df?|\xbab6\x00\x01\x06|\xfadSd\x00
\x00\x00\x00:\x0b\x00\x00\x00\x00\x00\x10\x11\x12\x13\x14\x15\x16\x17\x1
8\x19\x1a\x1b\x1c\x1d\x1e\x1f !"#$%&'()*+,-./01234567'

>>>
```

The MAC address in these two packets are not the same because we re-built the network architecture to change the tunnel mode, and since we were not interested in specifying static MAC address for `routeA` and `routeB`, they are assigned with another random MAC address after the network is re-built.

18

We can also try to retrieve the original packet from the payload of the encapsulated packets. For those in IP encapsulation mode, the original packet content is located after the 8<sup>th</sup> byte of the payload, while in UDP encapsulation mode, it is located after the 12<sup>th</sup> byte:

```

bell@Bell-Kirato: ~
[2/251]
>>> ip_p = rdpcap('encap_ip.pcap')
>>> Ether(ip_p[0].payload.load[8:]).show()
###[ Ethernet ]###
  dst      = 00:11:22:00:01:00
  src      = 00:11:22:00:03:00
  type     = n_802_1Q
###[ 802.1Q ]###
  prio     = 0
  id       = 0
  vlan     = 100
  type     = IPv4
###[ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 84
  id       = 33317
  flags    = DF
  frag     = 0
  ttl      = 64
  proto    = icmp
  chksum   = 0x6f2f
  src      = 192.168.100.2
  dst      = 192.168.100.1
  \options \
###[ ICMP ]###
  type     = echo-request
  code     = 0
  chksum   = 0xdf3f
  id       = 0xba36
  seq      = 0x1
  unused   = ''
###[ Raw ]###
  load     = '\x06\xffaSd\x00\x00\x00:W\x0b\x00\x00\x00\x00\x10\x11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !
"$%&'()*+,-./01234567'

[01 0:linux]
"Bell-Kirato" 21:18 04-May-21

[2/323]
>>> udp_p = rdpcap('encap_udp.pcap')
>>> Ether(udp_p[1].payload.load[12:]).show()
###[ Ethernet ]###
  dst      = 00:11:22:00:01:00
  src      = 00:11:22:00:03:00
  type     = n_802_1Q
###[ 802.1Q ]###
  prio     = 0
  id       = 0
  vlan     = 100
  type     = IPv4
###[ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 84
  id       = 24429
  flags    = 
  frag     = 0
  ttl      = 64
  proto    = icmp
  chksum   = 0xd1e7
  src      = 192.168.100.2
  dst      = 192.168.100.1
  \options \
###[ ICMP ]###
  type     = echo-reply
  code     = 0
  chksum   = 0xd38d
  id       = 0xb684
  seq      = 0x1
  unused   = ''
###[ Raw ]###
  load     = '\xff8Sd\x00\x00\x00:\xbd\x0c\x00\x00\x00\x10\x11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !
"$%&'()*+,-./01234567'

```

Figure 12: Obtaining original packet from payload in IP vs. UDP encap modes

## 5.2 Compare to tunnel made with GRE in GRETAP mode

Instead of L2TPv3 tunnel, it is also possible to set up GRE tunnel in GRETAP mode.

To configure rout1:

```

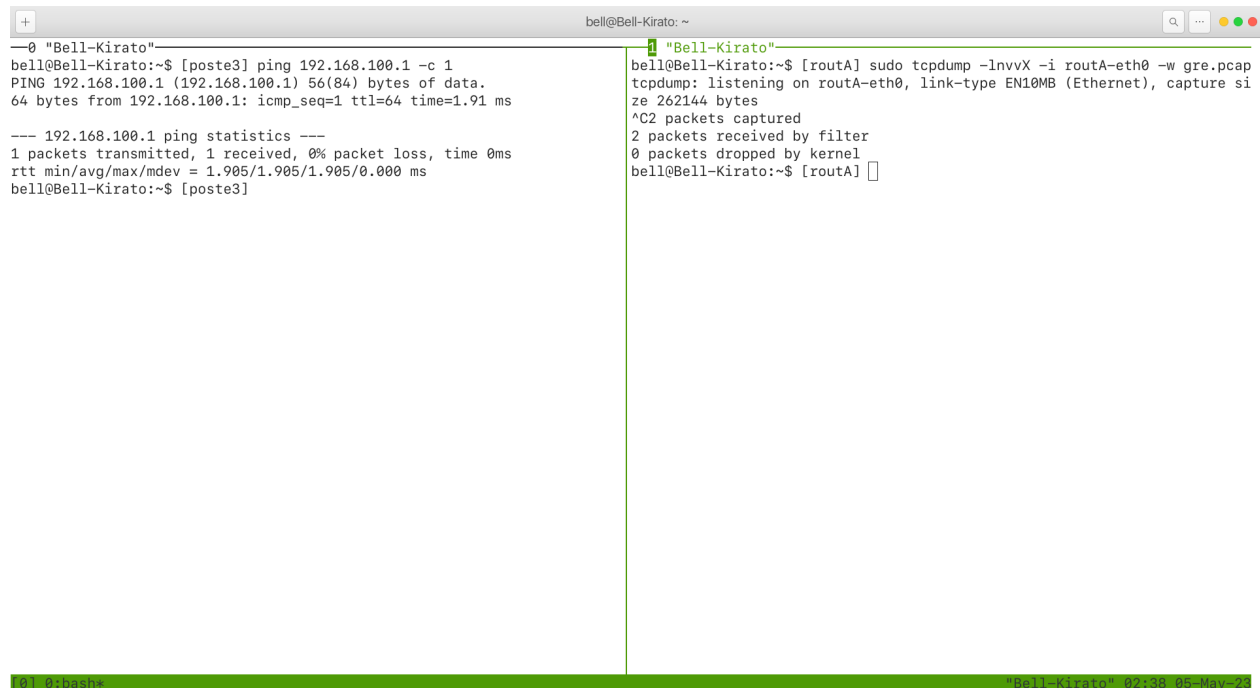
1 $ sudo ip netns exec rout1 ip l add eoip1 type gretap remote 172.16.2.253
   local 172.16.1.253 nopmtudisc
2 $ sudo ip netns exec rout1 ip l set dev eoip1 up
3 $ sudo ip netns exec rout1 brctl addbr tunnel
4 $ sudo ip netns exec rout1 brctl addif tunnel eoip1
5 $ sudo ip netns exec rout1 brctl addif tunnel rout1-eth1
6 $ sudo ip netns exec rout1 ip l set tunnel up
7 $ sudo ip netns exec rout1 ip l add link tunnel name tunnel.100 type vln
   id 100
8 $ sudo ip netns exec rout1 ip l set tunnel.100 up
9 $ sudo ip netns exec rout1 ip l add link tunnel name tunnel.200 type vln
   id 200
10 $ sudo ip netns exec rout1 ip l set tunnel.200 up
11 $ sudo ip netns exec rout1 ip a add 192.168.100.254/24 dev tunnel.100
12 $ sudo ip netns exec rout1 ip a add 192.168.200.254/24 dev tunnel.200

```

To configure rout2:

```
1 $ sudo ip netns exec rout2 ip 1 add eoip1 type gretap remote 172.16.1.253
   local 172.16.2.253 nopmtudisc
2 $ sudo ip netns exec rout2 ip 1 set dev eoip1 up
3 $ sudo ip netns exec rout2 brctl addbr tunnel
4 $ sudo ip netns exec rout2 brctl addif tunnel eoip1
5 $ sudo ip netns exec rout2 brctl addif tunnel rout2-eth1
6 $ sudo ip netns exec rout2 ip 1 set tunnel up
7 $ sudo ip netns exec rout2 ip 1 add link tunnel name tunnel.100 type vlan
   id 100
8 $ sudo ip netns exec rout2 ip 1 set tunnel.100 up
9 $ sudo ip netns exec rout2 ip 1 add link tunnel name tunnel.200 type vlan
   id 200
10 $ sudo ip netns exec rout2 ip 1 set tunnel.200 up
11 $ sudo ip netns exec rout2 ip a add 192.168.100.253/24 dev tunnel.100
12 $ sudo ip netns exec rout2 ip a add 192.168.200.253/24 dev tunnel.200
```

After having finished configuring the tunnel for rout1 and rout2, we can try to capture the traffic. In this example, once again we will capture the packets arrive at interface `routA-eth0` of `routA`, while sending ICMP requests from `poste3` to `poste1`:



```
bell@Bell-Kirato: ~  
--- "Bell-Kirato"  
bell@Bell-Kirato:~$ [poste3] ping 192.168.100.1 -c 1  
PING 192.168.100.1 (192.168.100.1) 56(84) bytes of data.  
64 bytes from 192.168.100.1: icmp_seq=1 ttl=64 time=1.91 ms  
  
--- 192.168.100.1 ping statistics ---  
1 packets transmitted, 1 received, 0% packet loss, time 0ms  
rtt min/avg/max/mdev = 1.905/1.905/1.905/0.000 ms  
bell@Bell-Kirato:~$ [poste3]  
  
bell@Bell-Kirato:~$ [routA] sudo tcpdump -lnvX -i routA-eth0 -w gre.pcap  
tcpdump: listening on routA-eth0, link-type EN10MB (Ethernet), capture si  
ze 262144 bytes  
^C2 packets captured  
2 packets received by filter  
0 packets dropped by kernel  
bell@Bell-Kirato:~$ [routA]
```

Figure 13: Capture traffic on `routA-eth0` after setting up GRE-TAP tunnel

Then, we can use `scapy` to analyze the difference between the packets encapsulated using L2TPv3 protocol and GRE/TAP protocol:

```

bell@Bell-Kirato: ~
--0 "Bell-Kirato"
>>> l2tp_p = rdpcap('encap_ip.pcap')
>>> l2tp_p[0].show()
#### [ Ethernet ]####
    dst      = c2:74:31:03:e6:8a
    src      = ae:8f:79:84:e4:fd
    type     = IPv4
#### [ IP ]####
    version  = 4
    ihl      = 5
    tos      = 0x0
    len      = 130
    id       = 11694
    flags    =
    frag     = 0
    ttl      = 64
    proto    = l2tp
    chksum   = 0xef40
    src      = 172.16.1.253
    dst      = 172.16.2.253
    \options \
#### [ Raw ]####
    load     = '\x00\x00\x07\x00\x00\x00\x00\x00\x00\x00\x11\x00\x01\x00\x03\x00\x81\x00\x0d\x00\x00E\x00\x00T\x82@\x00@\x01o/\xc0\xa8d\x02\xc0\xa8d\x01\x08\x00\xdf?\xbad\x00\x01\x06\xfaSd\x00\x00\x00\x00:W\x0b\x00\x00\x00\x00\x10\x11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !"#%&'()*+,-./01234567'
>>>

bell@Bell-Kirato: ~
>>> gre_p = rdpcap('gre.pcap')
>>> gre_p[0].show()
#### [ Ethernet ]####
    dst      = da:8f:ba:75:9a:02
    src      = 12:6d:53:23:ca:c5
    type     = IPv4
#### [ IP ]####
    version  = 4
    ihl      = 5
    tos      = 0x0
    len      = 126
    id       = 24284
    flags    =
    frag     = 0
    ttl      = 64
    proto    = gre
    chksum   = 0xbe5a
    src      = 172.16.1.253
    dst      = 172.16.2.253
    \options \
#### [ GRE ]####
    checksum_present= 0
    routing_present= 0
    key_present= 0
    seqnum_present= 0
    strict_route_source= 0
    recursion_control= 0
    flags         = 0
    version       = 0
    proto        = TEB
#### [ Ethernet ]####
    dst      = 00:11:22:00:01:00
    src      = 00:11:22:00:03:00
    type     = n_802_1Q
#### [ 802.1Q ]####

```

Figure 14: Packets encapsulated using L2TPv3 vs. GRE/TAP

Notice how the original packet remains its original form (with its own Ethernet header, 802.1Q or VLAN header, etc.) and are encapsulated with an additional GRE header, then another IP header and Ethernet header, instead of being encapsulated in the payload of the new packet in the form of raw bytes as in the case of L2TP protocol.

We can also see that the MTU on the tunnel interfaces of the routers are reduced to 1462, instead of 1500, due to the increasing information in packet headers.

```

1 $ sudo ip netns exec rout1 ip l show eoip1
2 5: eoip1@NONE: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1462 qdisc fq_codel
   master tunnel state UNKNOWN mode DEFAULT group default qlen 1000
3   link/ether ee:de:d5:f0:73:ea brd ff:ff:ff:ff:ff:ff

```

## 6 Setting up IPsec encryption on L2TPv3 tunnel

In order to set up IPsec encryption, combining both AH (*Authentication Header*) and ESP (*Encapsulating Security Payload*) in “transport” mode, on the L2TPv3 tunnel between `rout1` and `rout2`, we will utilize the `ip xfrm` command. Each protocol requires its own SA for each direction, so 4 SAs are needed to a VPN using AH+ESP.

First, we choose an arbitrary SPI (*Security Parameters Index*) for the SAs (*Security Associations*), and randomly generate the keys for SHA256 and AES protocols. AH SAs use 256-bit long and ESP SAs use 160-bit long keys:

```

1 # Choose a random SPI (Security Parameters Index)
2 index=0x12345678
3
4 # Generate random secret keys for symmetric encryption: AH SAs using 256
   bit (32 byte) long and ESP SAs using 160 bit (20 byte) long keys
5 ah_key_1=$(openssl rand -hex 32)
6 esp_key_1=$(openssl rand -hex 20)
7 ah_key_2=$(openssl rand -hex 32)
8 esp_key_2=$(openssl rand -hex 20)

```

After having the keys and SPI, we can configure the SAD (*Security Association Database*) using `ip xfrm state` and SPD (*Security Policy Database*) using `ip xfrm policy`. On `rout1`:

```

1 # # Config IPsec on rout1 end
2 # Flush the SAD and SPD
3 sudo ip netns exec rout1 ip xfrm state flush
4 sudo ip netns exec rout1 ip xfrm policy flush
5
6 # Security association
7 sudo ip netns exec rout1 ip xfrm state add src 172.16.1.253 dst
   172.16.2.253 proto esp spi $index reqid $index mode transport auth
   sha256 0x$ah_key_1 enc "rfc3686(ctr(aes))" 0x$esp_key_1
8 sudo ip netns exec rout1 ip xfrm state add src 172.16.2.253 dst
   172.16.1.253 proto esp spi $index reqid $index mode transport auth
   sha256 0x$ah_key_2 enc "rfc3686(ctr(aes))" 0x$esp_key_2
9
10 # Security policies
11 sudo ip netns exec rout1 ip xfrm policy add src 172.16.1.253 dst
   172.16.2.253 dir out tmpl src 172.16.1.253 dst 172.16.2.253 proto esp
   reqid $index mode transport
12 sudo ip netns exec rout1 ip xfrm policy add src 172.16.2.253 dst
   172.16.1.253 dir in tmpl src 172.16.2.253 dst 172.16.1.253 proto esp
   reqid $index mode transport

```

Symmetrically, on `rout2`:

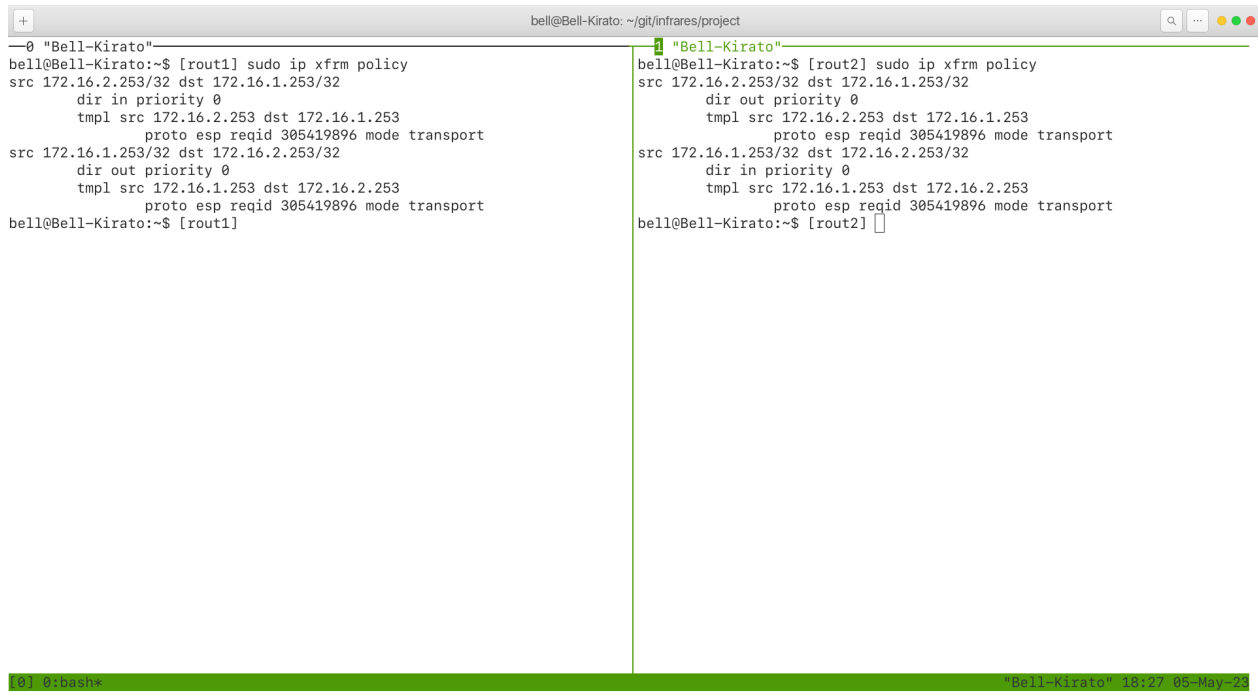
```

1 # # Config IPsec on rout2 end
2 # Flush the SAD and SPD
3 sudo ip netns exec rout2 ip xfrm state flush
4 sudo ip netns exec rout2 ip xfrm policy flush
5
6 # Security association
7 sudo ip netns exec rout2 ip xfrm state add src 172.16.1.253 dst
   172.16.2.253 proto esp spi $index reqid $index mode transport auth
   sha256 0x$ah_key_1 enc "rfc3686(ctr(aes))" 0x$esp_key_1
8 sudo ip netns exec rout2 ip xfrm state add src 172.16.2.253 dst
   172.16.1.253 proto esp spi $index reqid $index mode transport auth
   sha256 0x$ah_key_2 enc "rfc3686(ctr(aes))" 0x$esp_key_2
9
10 # Security policies
11 sudo ip netns exec rout2 ip xfrm policy add src 172.16.1.253 dst
   172.16.2.253 dir in tmpl src 172.16.1.253 dst 172.16.2.253 proto esp
   reqid $index mode transport
12 sudo ip netns exec rout2 ip xfrm policy add src 172.16.2.253 dst
   172.16.1.253 dir out tmpl src 172.16.2.253 dst 172.16.1.253 proto esp
   reqid $index mode transport

```

All the configurations are in the file `config_ipsec`.

We can check if the SAs are configured correctly using `ip xfrm policy` and `ip xfrm state`:

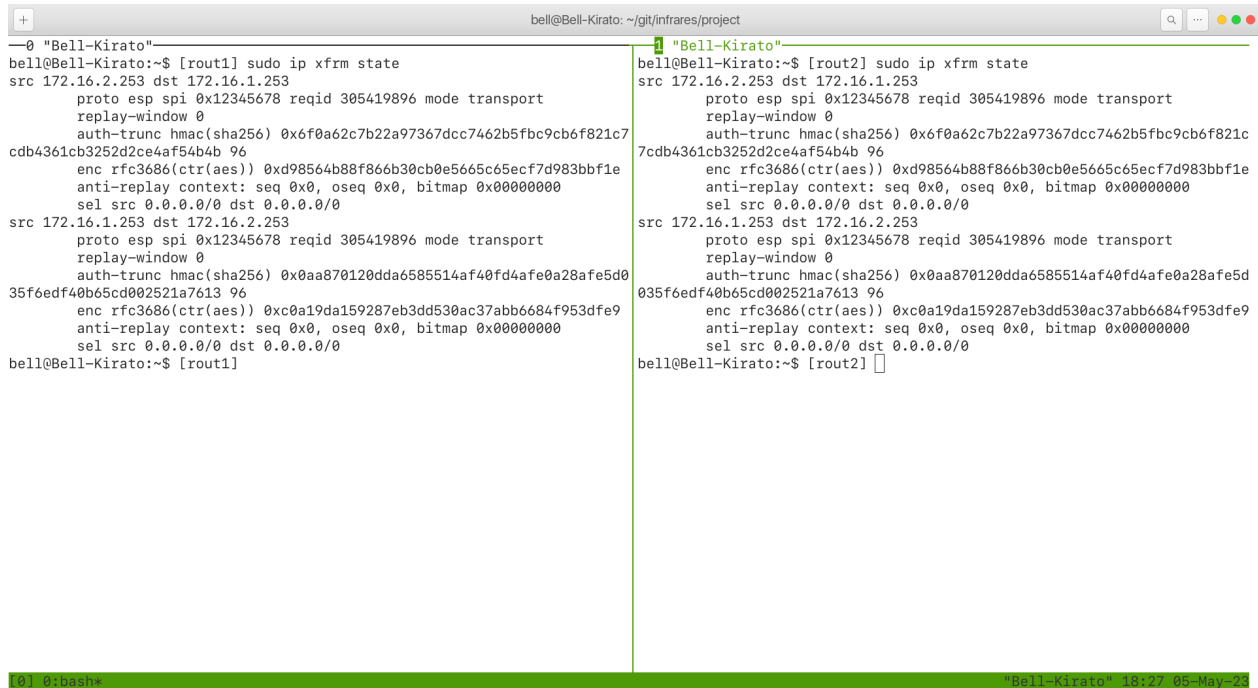


The image shows a terminal window with two panes. The left pane shows the output of `ip xfrm policy` on `route1`, and the right pane shows the output on `route2`. Both show a single policy for ESP mode transport with specific SPI and sequence number requirements.

```
bell@Bell-Kirato:~$ [route1] sudo ip xfrm policy
src 172.16.2.253/32 dst 172.16.1.253/32
  dir in priority 0
  tmpl src 172.16.2.253 dst 172.16.1.253
    proto esp reqid 305419896 mode transport
src 172.16.1.253/32 dst 172.16.2.253/32
  dir out priority 0
  tmpl src 172.16.1.253 dst 172.16.2.253
    proto esp reqid 305419896 mode transport
bell@Bell-Kirato:~$ [route1]
```

```
bell@Bell-Kirato:~$ [route2] sudo ip xfrm policy
src 172.16.2.253/32 dst 172.16.1.253/32
  dir out priority 0
  tmpl src 172.16.2.253 dst 172.16.1.253
    proto esp reqid 305419896 mode transport
src 172.16.1.253/32 dst 172.16.2.253/32
  dir in priority 0
  tmpl src 172.16.1.253 dst 172.16.2.253
    proto esp reqid 305419896 mode transport
bell@Bell-Kirato:~$ [route2]
```

Figure 15: Verify IPsec policy on `route1` and `route2`



The image shows a terminal window with two panes. The left pane shows the output of `ip xfrm state` on `route1`, and the right pane shows the output on `route2`. Both show the state of the ESP SAs, including SPI, HMAC, and replay window information.

```
bell@Bell-Kirato:~$ [route1] sudo ip xfrm state
src 172.16.2.253 dst 172.16.1.253
  proto esp spi 0x12345678 reqid 305419896 mode transport
  replay-window 0
  auth-trunc hmac(sha256) 0x6f0a62c7b22a97367dcc7462b5fbc9cb6f821c7
  enc rfc3686(ctr(aes)) 0xd98564b88f866b30cb0e5665c65ecf7d983bbf1e
  anti-replay context: seq 0x0, oseq 0x0, bitmap 0x00000000
  sel src 0.0.0.0/0 dst 0.0.0.0/0
src 172.16.1.253 dst 172.16.2.253
  proto esp spi 0x12345678 reqid 305419896 mode transport
  replay-window 0
  auth-trunc hmac(sha256) 0x0aa870120dda6585514af40fd4afe0a28afe5d0
  enc rfc3686(ctr(aes)) 0xc0a19da159287eb3dd530ac37abb6684f953dfe9
  anti-replay context: seq 0x0, oseq 0x0, bitmap 0x00000000
  sel src 0.0.0.0/0 dst 0.0.0.0/0
bell@Bell-Kirato:~$ [route1]
```

```
bell@Bell-Kirato:~$ [route2] sudo ip xfrm state
src 172.16.2.253 dst 172.16.1.253
  proto esp spi 0x12345678 reqid 305419896 mode transport
  replay-window 0
  auth-trunc hmac(sha256) 0x6f0a62c7b22a97367dcc7462b5fbc9cb6f821c
  enc rfc3686(ctr(aes)) 0xd98564b88f866b30cb0e5665c65ecf7d983bbf1e
  anti-replay context: seq 0x0, oseq 0x0, bitmap 0x00000000
  sel src 0.0.0.0/0 dst 0.0.0.0/0
src 172.16.1.253 dst 172.16.2.253
  proto esp spi 0x12345678 reqid 305419896 mode transport
  replay-window 0
  auth-trunc hmac(sha256) 0x0aa870120dda6585514af40fd4afe0a28afe5d
  enc rfc3686(ctr(aes)) 0xc0a19da159287eb3dd530ac37abb6684f953dfe9
  anti-replay context: seq 0x0, oseq 0x0, bitmap 0x00000000
  sel src 0.0.0.0/0 dst 0.0.0.0/0
bell@Bell-Kirato:~$ [route2]
```

Figure 16: Verify IPsec state on `route1` and `route2`



After verifying that IPsec is configured on **route1** and **route2** appropriately, the packets are captured on **routeA**'s interface **routeA-eth0** with **tcpdump**. We will try to ping from **postel** to **route1** and inspect the traffic:

```

bell@Bell-Kirato:~$ [routeA] sudo tcpdump -lnvv -i routeA-eth0
tcpdump: listening on routeA-eth0, link-type EN10MB (Ethernet), capture size 262144 bytes
18:33:29.271878 IP (tos 0x0, ttl 62, id 45977, offset 0, flags [none], protocol ESP (50), length 104)
172.16.2.253 > 172.16.1.253: ESP(spi=0x12345678,seq=0x3), length 84
18:33:29.272441 IP (tos 0x0, ttl 64, id 51960, offset 0, flags [none], protocol ESP (50), length 104)
172.16.1.253 > 172.16.2.253: ESP(spi=0x12345678,seq=0x2), length 84
18:33:29.273288 IP (tos 0x0, ttl 62, id 45978, offset 0, flags [none], protocol ESP (50), length 160)
172.16.2.253 > 172.16.1.253: ESP(spi=0x12345678,seq=0x4), length 140
18:33:29.273472 IP (tos 0x0, ttl 64, id 51961, offset 0, flags [none], protocol ESP (50), length 160)
172.16.1.253 > 172.16.2.253: ESP(spi=0x12345678,seq=0x3), length 140

bell@Bell-Kirato:~$ [postel] ping -c 1 192.168.100.254
PING 192.168.100.254 (192.168.100.254) 56(84) bytes of data:
64 bytes from 192.168.100.254: icmp_seq=1 ttl=64 time=2.63 ms

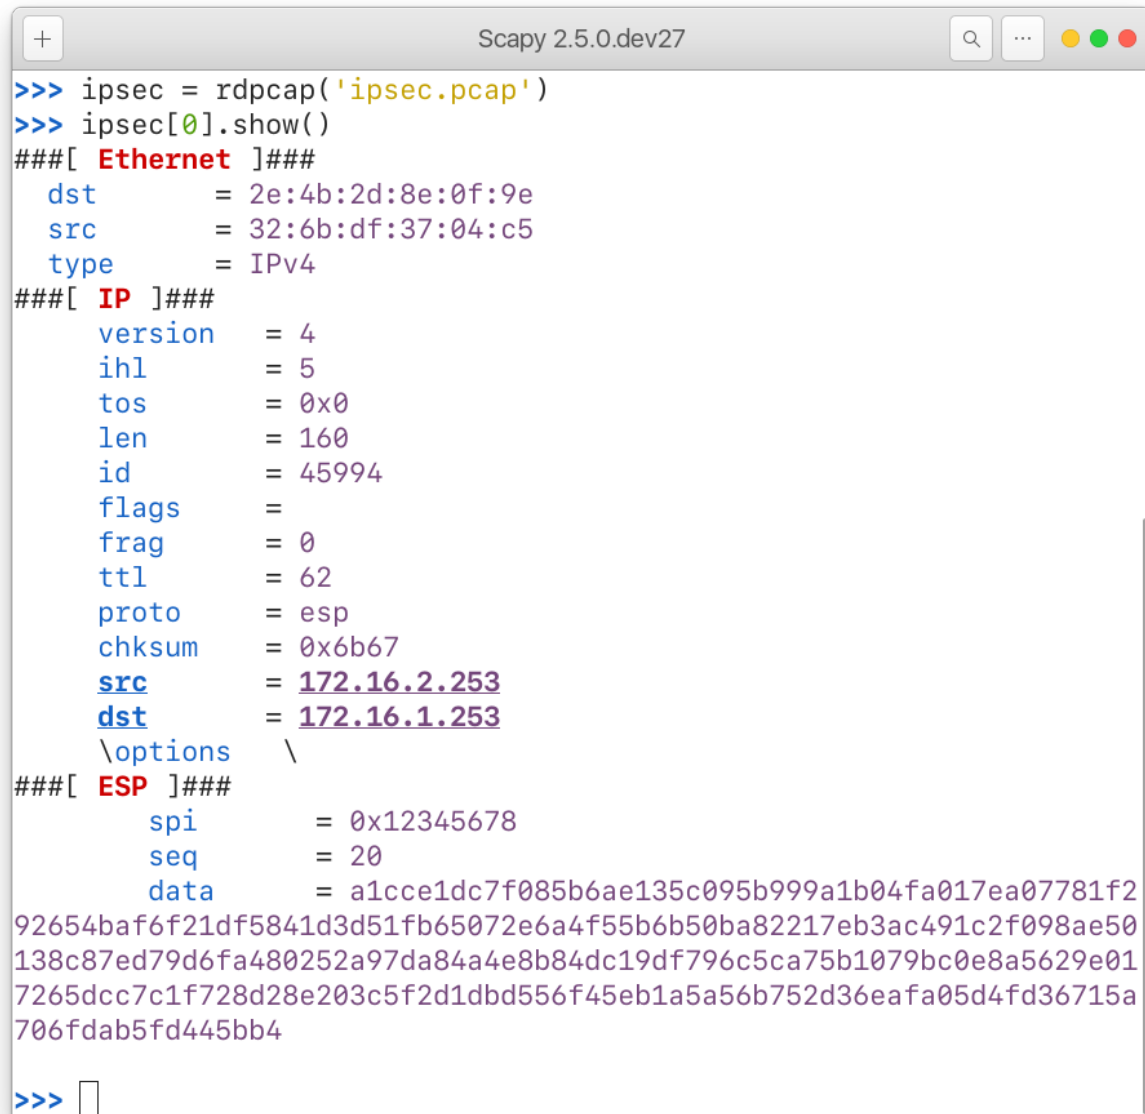
--- 192.168.100.254 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 2.628/2.628/2.628/0.000 ms
bell@Bell-Kirato:~$ [postel]

```

Figure 17: Inspect traffic going through **routeA** after implementing IPsec on L2TPv3 tunnel

It can be seen from Figure 25 that the packets are now ESP packets, which were encrypted with ESP protocol, instead of **ip-proto-115** packets as we have seen before.

We cannot access the content of the encrypted packets captured on an intermediary (routeA in this case) even if we try to use `scapy`, as the content of original packet was encrypted and can be seen from the field `data` of the ESP header:



```
>>> ipsec = rdpcap('ipsec.pcap')
>>> ipsec[0].show()
###[ Ethernet ]###
  dst      = 2e:4b:2d:8e:0f:9e
  src      = 32:6b:df:37:04:c5
  type     = IPv4
###[ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 160
  id       = 45994
  flags    =
  frag     = 0
  ttl      = 62
  proto    = esp
  checksum = 0x6b67
  src      = 172.16.2.253
  dst      = 172.16.1.253
  \options \
###[ ESP ]###
  spi      = 0x12345678
  seq      = 20
  data     = a1cce1dc7f085b6ae135c095b999a1b04fa017ea07781f2
92654baf6f21df5841d3d51fb65072e6a4f55b6b50ba82217eb3ac491c2f098ae50
138c87ed79d6fa480252a97da84a4e8b84dc19df796c5ca75b1079bc0e8a5629e01
7265dcc7c1f728d28e203c5f2d1dbd556f45eb1a5a56b752d36eafa05d4fd36715a
706fdab5fd445bb4
>>> 
```

Figure 18: Analyze the encrypted packet captured on an intermediary

The bandwidth between `poste1` and `rout1` sharply decreased after applying IPsec, due to the encrypting process:

```

bell@Bell-Kirato: ~/git/infrases/project
0 "Bell-Kirato"
[sudo] password for bell:
bell@Bell-Kirato:~/git/infrases/project$ [poste1] cd ~
bell@Bell-Kirato:~$ [poste1]
bell@Bell-Kirato:~$ [poste1] sudo iperf -c 192.168.100.254
Client connecting to 192.168.100.254, TCP port 5001
TCP window size: 926 KByte (default)
[ 3] local 192.168.100.1 port 49488 connected with 192.168.100.254 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.0 sec  1.09 GBytes  935 Mbits/sec
bell@Bell-Kirato:~$ [poste1]

1 "Bell-Kirato"
bell@Bell-Kirato:~$ [rout1]
bell@Bell-Kirato:~$ [rout1] sudo iperf -s
Server listening on TCP port 5001
TCP window size: 128 KByte (default)
[ 4] local 192.168.100.254 port 5001 connected with 192.168.100.1
port 49488
[ ID] Interval      Transfer    Bandwidth
[ 4] 0.0-10.0 sec  1.09 GBytes  932 Mbits/sec
^Cbell@Bell-Kirato:~$ [rout1]

```

Figure 19: Transfer speed and bandwidth before implementing IPsec

```

bell@Bell-Kirato: ~/git/infrases/project
0 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste1] sudo iperf -c 192.168.100.254
Client connecting to 192.168.100.254, TCP port 5001
TCP window size: 425 KByte (default)
[ 3] local 192.168.100.1 port 44236 connected with 192.168.100.254 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.0 sec  343 MBytes  287 Mbits/sec
bell@Bell-Kirato:~$ [poste1]

1 "Bell-Kirato"
bell@Bell-Kirato:~$ [rout1] sudo iperf -s
Server listening on TCP port 5001
TCP window size: 128 KByte (default)
[ 4] local 192.168.100.254 port 5001 connected with 192.168.100.1
port 44236
[ ID] Interval      Transfer    Bandwidth
[ 4] 0.0-10.1 sec  343 MBytes  285 Mbits/sec

```

Figure 20: Transfer speed and bandwidth after implementing IPsec

## 7 “Smart” Internet Access

### 7.1 Provide Internet access to the local network

To connect our virtual network to the real Internet, an interface `veth0` on the real Linux machine is created to connect itself to the switch `internet`, thus connect our local virtual network to the real Linux machine. We also need to assign an IP address to the interface `veth0`, which we choose to be `10.87.0.254`.

```
1 # Provide Internet access through the internet switch
2 ip l add veth0 type veth peer name internet-real
3 ovs-vsctl add-port internet internet-real
4 ip l set dev internet-real up
5 ip l set dev veth0 up
6 ip a add dev veth0 10.87.0.254/24
```

Then we modify the routing tables of the real Linux machine and the routers to enable the traffic between them, and we need to enable IPv4 forwarding on the Linux machine as well:

```
1 # Config the routing tables for the routers
2 ip netns exec routA ip r add default via 10.87.0.254 dev routA-eth1
3 ip netns exec rout1 ip r add default via 172.16.1.254 dev rout1-eth0
4 ip netns exec routB ip r add default via 10.87.0.254 dev routB-eth1
5 ip netns exec rout2 ip r add default via 172.16.2.254 dev rout2-eth0
6
7 # Config the routing table of the Linux machine
8 sysctl net.ipv4.conf.all.forwarding=1
9 ip r add 172.16.1.0/24 via 10.87.0.1 dev veth0
10 ip r add 172.16.2.0/24 via 10.87.0.2 dev veth0
```

Finally, we add the NetFilter rules to MASQUERADE all outgoing packets from the real Linux machine’s network interface `wlo1`, which is connected to the internet, since the local network is private, if MASQUERADE is not used, the Internet would not know how to return packets to IP addresses coming from it.

```
1 iptables -t nat -A POSTROUTING -o wlo1 -j MASQUERADE
```

Now, we can try to send a ping to the internet from the local network. Let's say we want to ping 8.8.8.8 from `router1`:

```

bell@Bell-Kirato: ~/git/infrases/project
--- "Bell-Kirato" ---
bell@Bell-Kirato:~$ [router1] ping -c 1 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=118 time=15.9 ms

--- 8.8.8.8 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 15.850/15.850/15.850/0.000 ms
bell@Bell-Kirato:~$ [router1]

1 "Bell-Kirato"
bell@Bell-Kirato:~$ sudo tcpdump -lnvv -i veth0 icmp
tcpdump: listening on veth0, link-type EN10MB (Ethernet), capture size 2
62144 bytes
03:49:27.548293 IP (tos 0x0, ttl 63, id 13626, offset 0, flags [DF], pro
to ICMP (1), length 84)
172.16.1.253 > 8.8.8.8: ICMP echo request, id 10702, seq 1, length 6
4
03:49:27.562895 IP (tos 0x0, ttl 119, id 0, offset 0, flags [none], prot
o ICMP (1), length 84)
8.8.8.8 > 172.16.1.253: ICMP echo reply, id 10702, seq 1, length 64

2 "Bell-Kirato"
bell@Bell-Kirato:~$ sudo tcpdump -lnvv -i wlo1 icmp
tcpdump: listening on wlo1, link-type EN10MB (Ethernet), capture size 26
2144 bytes
03:49:27.548333 IP (tos 0x0, ttl 62, id 13626, offset 0, flags [DF], pro
to ICMP (1), length 84)
10.192.131.215 > 8.8.8.8: ICMP echo request, id 10702, seq 1, length
64
03:49:27.562848 IP (tos 0x0, ttl 120, id 0, offset 0, flags [none], prot
o ICMP (1), length 84)
8.8.8.8 > 10.192.131.215: ICMP echo reply, id 10702, seq 1, length 6
4

```

Figure 21: Ping from `router1` to 8.8.8.8

We can see that the traffic is well-transferred. The MASQUERADE also functions as expected, as we sniff on both interfaces `veth0` (connecting the real Linux machine to the local virtual network) and `wlo1` (connecting the real Linux machine to the Internet).

## 7.2 Provide Internet access to the hosts

To provide Internet to the hosts from both VLANs: `poste1`, `poste2`, `poste3`, and `poste4`, we need to not only MASQUERADE the packets coming out of our real Linux machine, but also those coming out of `router1` and `router2`, to go to the `internet` switch as well. The reason is the same as before: the Linux machine does not know how to route packets to our hosts yet, so either we add specific routes to the the routing table of the real Linux machine to route packets directly to those hosts, or we MASQUERADE those coming out of `router1` and `router2`. Here the second solution is chosen.

```

1 ip netns exec rout1 iptables -t nat -A POSTROUTING -o rout1-eth0 -j
  MASQUERADE
2 ip netns exec rout2 iptables -t nat -A POSTROUTING -o rout2-eth0 -j
  MASQUERADE

```

The second line is actually unnecessary at this point, since the traffic from `poste1` and `poste2` would first be redirected through the tunnel to `router1`, then goes to the Internet. In brief, the right-hand side network will communicate with the Internet through `router1`, so there is no need to MASQUERADE traffic coming out of interface `rout2-eth0` of `router2` at this point. But it is still added as we are going to use it in the upcoming part of the exercise.

We can verify that the communication between `poste1` and the Internet is well-passed through `route1`:

```

bell@Bell-Kirato: ~/git/infrases/project

0 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste1] ping -c 1 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=117 time=17.6 ms

--- 8.8.8.8 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 17.568/17.568/17.568/0.000 ms
bell@Bell-Kirato:~$ [poste1]

1 "Bell-Kirato"
bell@Bell-Kirato:~$ [route1] sudo tcpdump -lnvv -i tunnel.100 icmp
tcpdump: listening on tunnel.100, link-type EN10MB (Ethernet), capture s
size 262144 bytes
04:23:25.825146 IP (tos 0x0, ttl 64, id 36280, offset 0, flags [DF], pro
to ICMP (1), length 84)
192.168.100.1 > 8.8.8.8: ICMP echo request, id 12275, seq 1, length
64
04:23:25.840964 IP (tos 0x0, ttl 117, id 0, offset 0, flags [none], prot
o ICMP (1), length 84)
8.8.8.8 > 192.168.100.1: ICMP echo reply, id 12275, seq 1, length 64

2 "Bell-Kirato"
bell@Bell-Kirato:~$ [route1] sudo tcpdump -lnvv -i route1-eth0 113/249
tcpdump: listening on route1-eth0, link-type EN10MB (Ethernet), capture s
size 262144 bytes
04:23:25.825146 IP (tos 0x0, ttl 62, id 29778, offset 0, flags [none], p
roto ESP (50), length 160)
172.16.2.253 > 172.16.1.253: ESP(spi=0x12345678,seq=0x1b), length 14
0
04:23:25.825347 IP (tos 0x0, ttl 63, id 36280, offset 0, flags [DF], pro
to ICMP (1), length 84)
172.16.1.253 > 8.8.8.8: ICMP echo request, id 12275, seq 1, length 6
4
04:23:25.840934 IP (tos 0x0, ttl 118, id 0, offset 0, flags [none], prot
o ICMP (1), length 84)
8.8.8.8 > 172.16.1.253: ICMP echo reply, id 12275, seq 1, length 64
04:23:25.841041 IP (tos 0x0, ttl 64, id 48330, offset 0, flags [none],

```

Figure 22: Ping from `poste1` to 8.8.8.8

### 7.3 Configure default gateway for the hosts

In this part, we have to provide the hosts with proper configuration using the DHCP server, so that the communication of `poste1` and `poste2` goes through `route2` instead of `route1` to optimize the Internet access.

First of all, this is where the MASQUERADE traffic from `route2` becomes useful, as previously mentioned in Subsection 7.2, since if we are going to communicate to the Internet directly through `route2`, it is necessary to MASQUERADE traffic coming out of `route2-eth0`, for the same reason as in the case of `route1`.

Now, to re-configure the hosts using DHCP protocol, we need to flush the DHCP cache memory on each of the hosts, and delete their default gateway in the routing table.

```

1 $ sudo ip netns exec poste1 dhclient -r -v
2 $ sudo ip netns exec poste1 dhclient -r poste1-eth0
3 $ sudo ip netns exec poste1 ip r del default
4 $ sudo ip netns exec poste2 dhclient -r -v
5 $ sudo ip netns exec poste2 dhclient -r poste2-eth0
6 $ sudo ip netns exec poste2 ip r del default
7 $ sudo ip netns exec poste3 dhclient -r -v
8 $ sudo ip netns exec poste3 dhclient -r poste3-eth0
9 $ sudo ip netns exec poste3 ip r del default
10 $ sudo ip netns exec poste4 dhclient -r -v
11 $ sudo ip netns exec poste4 dhclient -r poste4-eth0
12 $ sudo ip netns exec poste4 ip r del default

```

We update the configuration of the DHCP server that is set up on `route1` by modifying the file `/etc/dnsmasq.conf`:

```
1 dhcp-range=192.168.100.1,192.168.100.2,12h
2 dhcp-range=192.168.200.1,192.168.200.2,12h
3 dhcp-option=tag:route2.100,option:router,192.168.100.253
4 dhcp-option=tag:route2.200,option:router,192.168.200.253
5 dhcp-option=tag:route1.100,option:router,192.168.100.254
6 dhcp-option=tag:route1.200,option:router,192.168.200.254
7 dhcp-host=00:11:22:00:01:00, set:route2.100,192.168.100.1,12h
8 dhcp-host=00:11:22:00:02:00, set:route2.200,192.168.200.1,12h
9 dhcp-host=00:11:22:00:03:00, set:route1.100,192.168.100.2,12h
10 dhcp-host=00:11:22:00:04:00, set:route1.200,192.168.200.2,12h
```

The `dhcp-option` and `set` (in `dhcp-host`) are used to specify the default gateway for the hosts. Then, again, on `route1`, we start the DHCP server:

```
1 $ [route1] sudo dnsmasq -d -z
```

and on each host, we run the `dhclient`:

```
1 $ [<poste-name>] sudo dhclient -v <poste-name>-eth0
```

```
bell@Bell-Kirato: ~/git/infraes/project

-1 "Bell-Kirato"
bell@Bell-Kirato:~$ [route1] sudo dnsmasq -d -z
Internet Systems Consortium DHCP Client 4.4.1
Copyright 2004-2018 Internet Systems Consortium.
All rights reserved.
For info, please visit https://www.isc.org/software/dhcp/

Listening on LPF/poste1-eth0/00:11:22:00:01:00
Sending on LPF/poste1-eth0/00:11:22:00:01:00
Sending on Socket/fallback
DHCPDISCOVER on poste1-eth0 to 255.255.255.255 port 67 interval 3
(xid=0xc89af03)
DHCPOFFER of 192.168.100.1 from 192.168.100.254
DHCPREQUEST for 192.168.100.1 on poste1-eth0 to 255.255.255.255 port 67 (xid=0x3af890c)
DHCPACK of 192.168.100.1 from 192.168.100.254 (xid=0xc89af03)
bound to 192.168.100.1 -- renewal in 20994 seconds.
bell@Bell-Kirato:~$ [poste1]

-2 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste3] sudo dhclient -v poste3-eth0
Internet Systems Consortium DHCP Client 4.4.1
Copyright 2004-2018 Internet Systems Consortium.
All rights reserved.
For info, please visit https://www.isc.org/software/dhcp/

Listening on LPF/poste3-eth0/00:11:22:00:03:00
Sending on LPF/poste3-eth0/00:11:22:00:03:00
Sending on Socket/fallback
DHCPDISCOVER on poste3-eth0 to 255.255.255.255 port 67 interval 3
(xid=0x4b21fb3d)
DHCPOFFER of 192.168.100.2 from 192.168.100.254
DHCPREQUEST for 192.168.100.2 on poste3-eth0 to 255.255.255.255 port 67 (xid=0x3dffb214b)
DHCPACK of 192.168.100.2 from 192.168.100.254 (xid=0x4b21fb3d)
bound to 192.168.100.2 -- renewal in 17885 seconds.
bell@Bell-Kirato:~$ [poste3]

-3 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste2] sudo dhclient -v poste2-eth0
Internet Systems Consortium DHCP Client 4.4.1
Copyright 2004-2018 Internet Systems Consortium.
All rights reserved.
For info, please visit https://www.isc.org/software/dhcp/

Listening on LPF/poste2-eth0/00:11:22:00:02:00
Sending on LPF/poste2-eth0/00:11:22:00:02:00
Sending on Socket/fallback
DHCPDISCOVER on poste2-eth0 to 255.255.255.255 port 67 interval 3
(xid=0x4f871102)
DHCPOFFER of 192.168.200.1 from 192.168.200.254
DHCPREQUEST for 192.168.200.1 on poste2-eth0 to 255.255.255.255 port 67 (xid=0x211874f)
DHCPACK of 192.168.200.1 from 192.168.200.254 (xid=0x4f871102)
bound to 192.168.200.1 -- renewal in 16956 seconds.
bell@Bell-Kirato:~$ [poste2]

-4 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste4] sudo dhclient -v poste4-eth0
Internet Systems Consortium DHCP Client 4.4.1
Copyright 2004-2018 Internet Systems Consortium.
All rights reserved.
For info, please visit https://www.isc.org/software/dhcp/

Listening on LPF/poste4-eth0/00:11:22:00:04:00
Sending on LPF/poste4-eth0/00:11:22:00:04:00
Sending on Socket/fallback
DHCPDISCOVER on poste4-eth0 to 255.255.255.255 port 67 interval 3
(xid=0xc2f45431)
DHCPOFFER of 192.168.200.2 from 192.168.200.254
DHCPREQUEST for 192.168.200.2 on poste4-eth0 to 255.255.255.255 port 67 (xid=0x3154f4c2)
DHCPACK of 192.168.200.2 from 192.168.200.254 (xid=0xc2f45431)
bound to 192.168.200.2 -- renewal in 21294 seconds.
bell@Bell-Kirato:~$ [poste4]
```

Figure 23: The DHCP Protocol captured on `route1` and `poste1`, `poste2`, `poste3`, and `poste4` (with default gateway config by DHCP)

We can verify again that the default route on each of the hosts has changed:

```

bell@Bell-Kirato: ~/git/infraes/project
--0 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste1] ip r show default
default via 192.168.100.253 dev poste1-eth0
bell@Bell-Kirato:~$ [poste1]

--2 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste2] ip r show default
default via 192.168.200.253 dev poste2-eth0
bell@Bell-Kirato:~$ [poste2]

--1 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste3] ip r show default
default via 192.168.100.254 dev poste3-eth0
bell@Bell-Kirato:~$ [poste3]

--3 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste4] ip r show default
default via 192.168.200.254 dev poste4-eth0
bell@Bell-Kirato:~$ [poste4]

```

Figure 24: Modified default route on poste1, poste2, poste3, and poste4

And we verify that the traffic from poste1 to the Internet and back pass through rout2 instead of rout1:

```

bell@Bell-Kirato: ~/git/infraes/project
--0 "Bell-Kirato"
bell@Bell-Kirato:~$ [poste1] ping 8.8.8.8 -c 1
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data:
64 bytes from 8.8.8.8: icmp_seq=1 ttl=117 time
=16.2 ms

--- 8.8.8.8 ping statistics ---
1 packets transmitted, 1 received, 0% packet l
oss, time 0ms
rtt min/avg/max/mdev = 16.211/16.211/16.211/0.
000 ms
bell@Bell-Kirato:~$ [poste1]

--1 "Bell-Kirato"
bell@Bell-Kirato:~$ [rout1] sudo tcpdump -lnvv -i rout1-eth0 icmp
tcpdump: listening on rout1-eth0, link-type EN10MB (Ethernet), ca
pture size 262144 bytes

--2 "Bell-Kirato"
bell@Bell-Kirato:~$ [rout2] sudo tcpdump -lnvv -i rout2-eth0
icmp
tcpdump: listening on rout2-eth0, link-type EN10MB (Ethernet
), capture size 262144 bytes
05:12:31.910093 IP (tos 0x0, ttl 63, id 60966, offset 0, fla
gs [DF], proto ICMP (1), length 84)
172.16.2.253 > 8.8.8.8: ICMP echo request, id 15203, seq
1, length 64
05:12:31.925758 IP (tos 0x0, ttl 118, id 0, offset 0, flags
[none], proto ICMP (1), length 84)
8.8.8.8 > 172.16.2.253: ICMP echo reply, id 15203, seq 1
, length 64

```

Figure 25: Communication between poste1 and the Internet pass through rout2



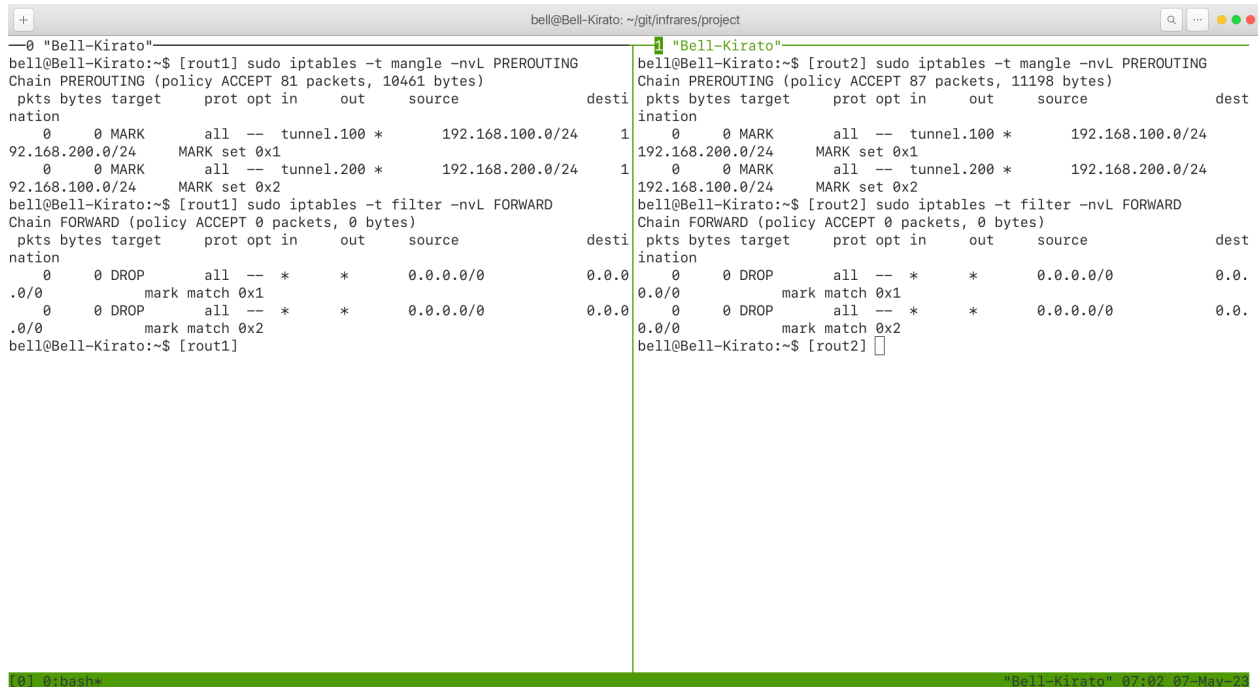
## 8 Blocking traffic between VLANs

Theoretically, traffics between different VLANs should be isolated (which is exactly the purpose of using VLAN), so we would want to block the traffic between them. We can do this either by using NetFilter or by using Routing Policy.

### 8.1 Using iptables rules

To isolate the traffic between VLAN100 and VLAN200, we would simply want to filter out (drop) any packets having source IP address from one VLAN and destination IP address from the other. These packets, indeed, should already have been decrypted, decapsulated from any kind of encapsulation, etc. Therefore, the configuration on **route1** and **route2** should look like this:

```
1 ip netns exec rout1 iptables -t mangle -A PREROUTING -i tunnel.100 -s
  192.168.100.0/24 -d 192.168.200.0/24 -j MARK --set-mark 1
2 ip netns exec rout1 iptables -t mangle -A PREROUTING -i tunnel.200 -s
  192.168.200.0/24 -d 192.168.100.0/24 -j MARK --set-mark 2
3 ip netns exec rout1 iptables -t filter -A FORWARD -m mark --mark 1 -j DROP
4 ip netns exec rout1 iptables -t filter -A FORWARD -m mark --mark 2 -j DROP
5
6 ip netns exec rout2 iptables -t mangle -A PREROUTING -i tunnel.100 -s
  192.168.100.0/24 -d 192.168.200.0/24 -j MARK --set-mark 1
7 ip netns exec rout2 iptables -t mangle -A PREROUTING -i tunnel.200 -s
  192.168.200.0/24 -d 192.168.100.0/24 -j MARK --set-mark 2
8 ip netns exec rout2 iptables -t filter -A FORWARD -m mark --mark 1 -j DROP
9 ip netns exec rout2 iptables -t filter -A FORWARD -m mark --mark 2 -j DROP
```



```
+
bell@Bell-Kirato: ~/git/infraes/project

0 "Bell-Kirato"
bell@Bell-Kirato:~$ [rout1] sudo iptables -t mangle -nvl PREROUTING
Chain PREROUTING (policy ACCEPT 81 packets, 10461 bytes)
pkts bytes target      prot opt in      out     source      desti
nation
0      0 MARK      all  --  tunnel.100 *      192.168.100.0/24
92.168.200.0/24 MARK set 0x1
0      0 MARK      all  --  tunnel.200 *      192.168.200.0/24
92.168.100.0/24 MARK set 0x2
bell@Bell-Kirato:~$ [rout1] sudo iptables -t filter -nvl FORWARD
Chain FORWARD (policy ACCEPT 0 packets, 0 bytes)
pkts bytes target      prot opt in      out     source      desti
nation
0      0 DROP      all  --  *      *      0.0.0.0/0   0.0.0
.0/0      mark match 0x1
0      0 DROP      all  --  *      *      0.0.0.0/0   0.0.0
.0/0      mark match 0x2
bell@Bell-Kirato:~$ [rout1]

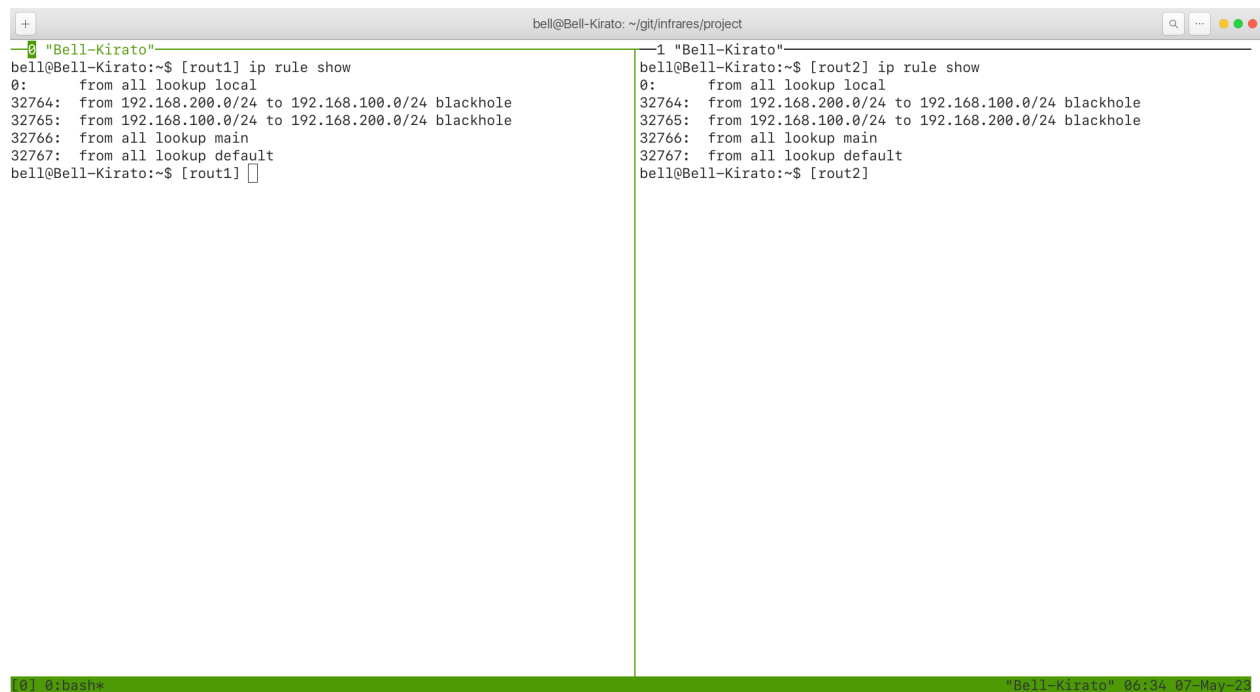
1 "Bell-Kirato"
bell@Bell-Kirato:~$ [rout2] sudo iptables -t mangle -nvl PREROUTING
Chain PREROUTING (policy ACCEPT 87 packets, 11198 bytes)
pkts bytes target      prot opt in      out     source      desti
ination
0      0 MARK      all  --  tunnel.100 *      192.168.100.0/24
192.168.200.0/24 MARK set 0x1
0      0 MARK      all  --  tunnel.200 *      192.168.200.0/24
192.168.100.0/24 MARK set 0x2
bell@Bell-Kirato:~$ [rout2] sudo iptables -t filter -nvl FORWARD
Chain FORWARD (policy ACCEPT 0 packets, 0 bytes)
pkts bytes target      prot opt in      out     source      desti
ination
0      0 DROP      all  --  *      *      0.0.0.0/0   0.0.
0.0/0      mark match 0x1
0      0 DROP      all  --  *      *      0.0.0.0/0   0.0.
0.0/0      mark match 0x2
bell@Bell-Kirato:~$ [rout2]
```

Figure 26: NetFilter rules on **route1** and **route2**

## 8.2 Using Routing Policy

In this case, we will take advantage of the `blackhole` route in Routing Policy. So instead of dropping packets that are from one VLAN to another as in `iptables`, we create Routing Policies that put them in the `blackhole` route, using `ip rule`:

```
1 ip netns exec rout1 ip rule add from 192.168.100.0/24 to 192.168.200.0/24
  blackhole
2 ip netns exec rout1 ip rule add from 192.168.200.0/24 to 192.168.100.0/24
  blackhole
3 ip netns exec rout2 ip rule add from 192.168.100.0/24 to 192.168.200.0/24
  blackhole
4 ip netns exec rout2 ip rule add from 192.168.200.0/24 to 192.168.100.0/24
  blackhole
```



```
bell@Bell-Kirato: ~/git/infraes/project

0 "Bell-Kirato"
bell@Bell-Kirato:~$ [rout1] ip rule show
0:      from all lookup local
32764:  from 192.168.200.0/24 to 192.168.100.0/24 blackhole
32765:  from 192.168.100.0/24 to 192.168.200.0/24 blackhole
32766:  from all lookup main
32767:  from all lookup default
bell@Bell-Kirato:~$ [rout1]

1 "Bell-Kirato"
bell@Bell-Kirato:~$ [rout2] ip rule show
0:      from all lookup local
32764:  from 192.168.200.0/24 to 192.168.100.0/24 blackhole
32765:  from 192.168.100.0/24 to 192.168.200.0/24 blackhole
32766:  from all lookup main
32767:  from all lookup default
bell@Bell-Kirato:~$ [rout2]
```

Figure 27: Routing Policies on `rout1` and `rout2`